

NAIVE BAYES ALGORITHM FOR TWITTER SENTIMENT ANALYSIS AND ITS
IMPLEMENTATION IN MAPREDUCE

A Thesis
Presented to
The Faculty of the Graduate School
At the University of Missouri

In Partial Fulfillment
Of the Requirements for the Degree
Master of Science

By
ZHAOYU LI
Dr. Yi Shang, Advisor

DECEMBER 2014

The undersigned, appointed by the dean of the Graduate School, have examined the
thesis entitled

NAIVE BAYES ALGORITHM FOR TWITTER SENTIMENT ANALYSIS AND ITS
IMPLEMENTATION IN MAPREDUCE

Presented by Zhaoyu Li

A candidate for the degree of

Master of Science

And hereby certify that, in their opinion, it is worthy of acceptance.

Dr. Yi Shang

Dr. Dong Xu

Dr. Jeffrey Uhlmann

ACKNOWLEDGEMENTS

I would like to first thank my advisor, Dr. Yi Shang. He showed me how to do research in computer science, and he always supported and inspired me through the whole development of my thesis. He provided me many research ideas and helped me produce solid works.

I would also like to thank other lab mates in Dr. Shang's lab, both those who have left and those who have just begun. I am glad that I have known them and I really enjoyed their company along the way.

Finally I would like to thank my committee members Dr. Dong Xu and Dr. Jeffrey Uhlmann for their support on this thesis.

TABLE OF CONTENTS

Acknowledgements.....	ii
List of Figures.....	v
Abstract.....	vii
1 . INTRODUCTION	1
1.1 PROBLEM AND MOTIVATION.....	3
1.2 CONTRIBUTIONS	3
1.3 THESIS ORGANIZATION	4
2 . BACKGROUND AND RELATED WORK.....	6
2.1 TEXT CLASSIFICATION	6
2.1.1 <i>Definition of Text Classification</i>	6
2.1.2 <i>Process of Text Classification</i>	6
2.2 SENTIMENT ANALYSIS	8
2.2.1 <i>Algorithms</i>	9
2.2.2 <i>Twitter Sentiment Analysis and Stock Price</i>	13
2.3 CURRENT DISTRIBUTED DATA PROCESSING SYSTEM.....	14
2.4 MAPREDUCE MODEL	14
2.5 AN OVERVIEW OF HADOOP MAPREDUCE CLUSTER	16
2.5.1 <i>Hadoop Cluster Infrastructure</i>	17
2.5.2 <i>Hadoop Distributed File System (HDFS)</i>	18
2.5.3 <i>MapReduce</i>	18
3 . NAIVE BAYES FOR TWITTER SENTIMENT ANALYSIS	21

3.1 GEOGRAPHICAL SENTIMENT AND STOCK PRICE.....	21
3.2 THE SOURCE DATA.....	22
3.3 SYSTEM METHODS AND IMPLEMENTATION	23
3.3.1 <i>Three Layer Design</i>	23
3.3.2 <i>Stock Market Sector</i>	24
3.3.3 <i>Geographical Location</i>	25
3.4 EXPERIMENTAL RESULTS.....	26
3.4.1 <i>Result Representations</i>	27
3.4.2 <i>Analysis of the Results</i>	28
4 . NAIVE BAYES MAPREDUCE IMPLEMENTATION.....	31
4.1 METHODS OVERVIEW.....	31
4.2 DATA PRE-PROCESSING.....	32
4.3 PARALLEL NAIVE BAYES IN MAPREDUCE.....	35
4.3.1 <i>Algorithm Implementation</i>	35
4.3.2 <i>Algorithm Time Complexity</i>	43
5 . EXPERIMENTAL RESULTS	45
5.1 METRICS SUMMARY	45
5.1.1 <i>Terminologies</i>	45
5.1.2 <i>Impact of Node Number</i>	47
5.1.3 <i>Impact of Data Size</i>	47
5.1.4 <i>Impact of Global Counters</i>	47
5.2 EVALUATION SCENARIOS.....	48
5.2.1 <i>Dataset</i>	48
5.2.2 <i>Hardware Configuration</i>	48

5.3 CPU TIME	51
5.4 DATA SIZE	53
5.4.1 <i>Total Execution Time</i>	54
5.4.2 <i>Total Slot Execution Time</i>	56
5.5 NUMBER OF NODES	59
5.5.1 <i>Total Execution Time</i>	59
5.5.2 <i>Total Slot Execution Time</i>	61
5.5.3 <i>Speedup</i>	64
5.6 HARDWARE CONFIGURATION	67
5.7 GLOBAL COUNTERS.....	68
6 . CONCLUSION	71
7 . FUTURE WORK	72
8 . BIBLIOGRAPHY	74

LIST OF FIGURES

Figure 1: Supervised Text Classification	8
Figure 2: Support Vector Machine on Classification.....	12
Figure 3: Hadoop Cluster Infrastructure	17
Figure 4: Typical Components of a Hadoop Cluster.....	20
Figure 5: Three-Layer Architecture of the System	23
Figure 6: Sentiment of Starbucks Corporation by City.....	28
Figure 7: Three Days of Sentiment of Starbucks Corporation by City	29
Figure 8: Whole System Method Overview.....	32
Figure 9: Tweets Data Pre-processing	33
Figure 10: Basic Amazon EMR Infrastructure	50
Figure 11: CPU Time and Number of Nodes.....	52
Figure 12: CPU Time and Size of Data.....	53
Figure 13: The Total Execution Time and Size of Data	54
Figure 14: The Total Slot Time and Size of Input Data.....	56
Figure 15: The Difference Time and Size of Input Data.....	57
Figure 16: The Total Execution Time and The Number of Nodes	59
Figure 17: The Total Slot Execution Time and The Number of Nodes	61
Figure 18: The Difference Time and Size of Input Data.....	63
Figure 19: Speedup and Size of Data.....	65
Figure 20: Speedup and The Number of Nodes.....	66
Figure 21: Total Execution Time and Size of Data From Two Clusters.....	68
Figure 22: The Total Time Execution and The Size of Data About Counters.....	69

ABSTRACT

The scale of social network data generated and processed is increasing exponentially in the Big Data era. Sentiment analysis aims to determine the attitude of a speaker or a writer with respect to some topic or the overall contextual polarity of a document, and the sentiment analysis on Twitter has also been used as a valid indicator of stock prices in the past. Naive Bayes is an algorithm to perform sentiment analysis. MapReduce programming model provides a simple and powerful model to implement distributed applications without having deeper knowledge of parallel programming. When a new hypothetical MapReduce sentiment analysis system is built to provide certain performance goal, we are lack of the benchmark and the traditional trial-and-error solution is extremely time-consuming and costly.

In this thesis we implemented a prototype system using Naive Bayes to find the correlation between the geographical sentiment on Twitter and the stock price behavior of companies. Also we implemented the Naive Bayes sentiment analysis algorithm in MapReduce model based on Hadoop, and evaluated the algorithm on large amount of Twitter data with different metrics. Based on the evaluation results, we provided a comprehensive MapReduce performance prediction model for Naive Bayes based sentiment analysis algorithm. The prediction model can predict task execution performance within a window, and can also be used by other MapReduce systems as a benchmark in order to improve the performance.

1. INTRODUCTION

Data has been growing exponentially in recent years. With the development of information highway, data can be generated and collected very fast, and the data is so large that it has exceeded the limit of our conventional processing methods and applications. There are tons of data generated in many fields everyday, such as the travelling data from cars and airplanes, the transaction data from stock market and banking system, or the data we created in our daily life. We have entered into a Big Data era [1]. Big Data is not only about the data volume, but also about variety and velocity. It can be both structured and unstructured data, in a dumped file or in real-time streaming format with high velocity. Data is valuable, and it always contains information that may be useful in many aspects, such as performing analysis, making decisions, or making predictions. Since these properties of Big Data, the traditional data mining methods are limited to find the useful information so that many Big Data processing methods and applications appeared. MapReduce [2] is a parallel programming model that can run on large dataset with high scalability and efficiency. Hadoop is one of many frameworks to implement MapReduce model and it has been widely used for both industries and academic research.

The social network is one of many data explosion areas. In old days, people communicated with each other around them. The spread of information is limited into small circles. Nowadays, thanks to the development of Internet, people talked everything on social network, including publishing posts, uploading pictures and sharing interesting videos. The content on social network are all user-generated. There is a new word called

“Internet Minute”, which refers to what happens in a minute on the Internet. According to the report from Intel [5], 347,222 tweets were published on the Twitter; 138,889 hours of video were watched on the YouTube, and 38,194 pictures were uploaded on the Facebook, in a single Internet Minute. The social network has become a necessary part of our daily life, and the amount of data generated is still increasing.

Among all social network medias, Twitter has become one of the most important platforms to share and communicate with friends. People tweets about various topics on movies, products, brands, and many others. Also, Twitter allows people to publish only 140 characters in a single tweet, which makes the information easy to read and spread. The most important and valuable data in these tweets is the sentiment of the publishers. When people published a tweet, they had their own feelings and attitudes towards the thing they talked about, such as satisfaction or un-satisfaction, positive feeling or negative feeling. These sentiment data would be a great source for companies or institutions to do marketing research and customer survey. It has been discussed that there is a correlation between public opinions and the stock price [24]. One of the stock price indicators is the market, and customers’ behavior will have a significant impact on the market. So generally, the public sentiment about a company and its products is proportion to its stock price behavior.

What is more, the existing data mining technologies are not able to handle this large amount of data. The computation requirement has increased far beyond our current machines and algorithms for Big Data. So we need to find out how to implement and perform sentiment analysis on social network data, especially the Twitter, with Big Data technologies efficiently.

1.1 Problem and Motivation

Little work has been done to actually expand on the topic of the correlation between Twitter sentiment and stock price. There, for instance, is no time frame as to how long it takes for a company's stock to act in accordance to its associated twitter sentiment, or if there is any demographic on Twitter that does a slightly better job of predicting a stock's behavior than another. This thesis hopes to develop a more precise explanation - using sentiment analysis, sorted by geographic location - as to why Twitter can be used for stock prediction, and in what ways can we make data mining via twitter more efficient.

In the mean time, it is a good idea to use Big Data technologies to perform sentiment analysis. However, for many sentiment analysis algorithms, they are usually implemented in sequential. The traditional parallel methods, such as MPI [3], are lack of scalability and ease of use. MapReduce is a good choice for Big Data solution. However, we don't have benchmark, and the traditional trial-and-error solution is extremely time-consuming and costly. When we provide a hypothetical MapReduce system for sentiment analysis, we usually want to know how to choose the cluster, how to optimize the program, how many nodes do we need, or how to get the best performance in specific scenarios.

1.2 Contributions

This thesis makes the following contributions:

1. This thesis describes preliminary efforts towards a system that collects tweets about a specific company or product and filters all the tweets with

geographical locations information. In this system, we attempted to establish a methodology that can be used for tracking attitude of a particular company found on Twitter using Naive Bayes sentiment analysis algorithm and a stock's behavior.

2. Naive Bayes based sentiment analysis algorithm in MapReduce model was implemented successfully. In order to test the performance of an algorithm with different metrics, we first need a runnable program that can perform sentiment analysis on Hadoop [4] cluster. The parallel MapReduce algorithm implementation is adopted from sequential one and has been optimized specifically for MapReduce.
3. We provided a comprehensive MapReduce performance prediction model for Naive Bayes based sentiment analysis algorithm related to the size of input data, the number of nodes, the use of global counters, and the hardware configuration, etc. The prediction model can predict task execution performance within a window, and can be used by other MapReduce systems as a benchmark in order to improve the performance.

1.3 Thesis Organization

The rest of the thesis is organized as follows:

1. Chapter 2 is about the background and related work of Twitter sentiment analysis algorithms and MapReduce model.
2. Chapter 3 is about the system to find the correlation between the geographical Twitter sentiment and the stock price.

3. Chapter 4 is about the MapReduce implementation of Naive Bayes sentiment analysis algorithm, including the algorithm design, system framework, and evaluation methods.
4. Chapter 5 is about the experiments we performed, the dataset, the analysis of experimental results and explanations of the plotted graphs, and based on that we provided a comprehensive prediction model for Naive Bayes MapReduce model.
5. Chapter 6 is about the conclusion and discussion of our work in this thesis.
6. Chapter 7 is about the future work and further areas of this topic leading from our work.

2. BACKGROUND AND RELATED WORK

In this chapter, we first review related work of text classification, and sentiment analysis algorithms. Then we have an overview of the current distributed data processing systems, what the MapReduce model is in general, and the Hadoop cluster including cluster infrastructure, the Hadoop distributed file system, and the MapReduce in Hadoop. After that, we review the MapReduce performance monitoring and modeling of Hadoop.

2.1 Text Classification

2.1.1 Definition of Text Classification

Text classification is a part of data mining to classify text based on the content of it. We can classify news, articles or books into different categories according to some features we defined.

2.1.2 Process of Text Classification

When we talk about text classification, we usually talk about the supervised classification, which has two stages: the training stage and the testing stage. Usually the training stage includes creating the labeled corpora dataset, pre-processing the training text, vectorization of the text, and training of the classifier. The testing stage includes pre-processing of testing text, vectorization and classification of the testing text. This process is illustrated in Figure 1.

(1) Creating Corpus

Collection of text based on categories. Every text belongs to one category and has been correctly labeled. Sometimes we divided this corpus into two sets: the training set and the testing set.

(2) Pre-processing

Remove all the unnecessary elements in the text, such as stopword, punctuation, or unreadable text. This step is very important because it will affect the training of the classifier.

(3) Vectorization of Text

Transform the text into vector that can be recognized for computer. All text will be represented as feature vector based on the features we selected.

(4) Training of the Classifier

Choose one of the text classification algorithms and feed the training corpus to the classifier to get a training model.

(5) Classification

After we get the training model, we can feed the testing data into it and get the prediction of classification.

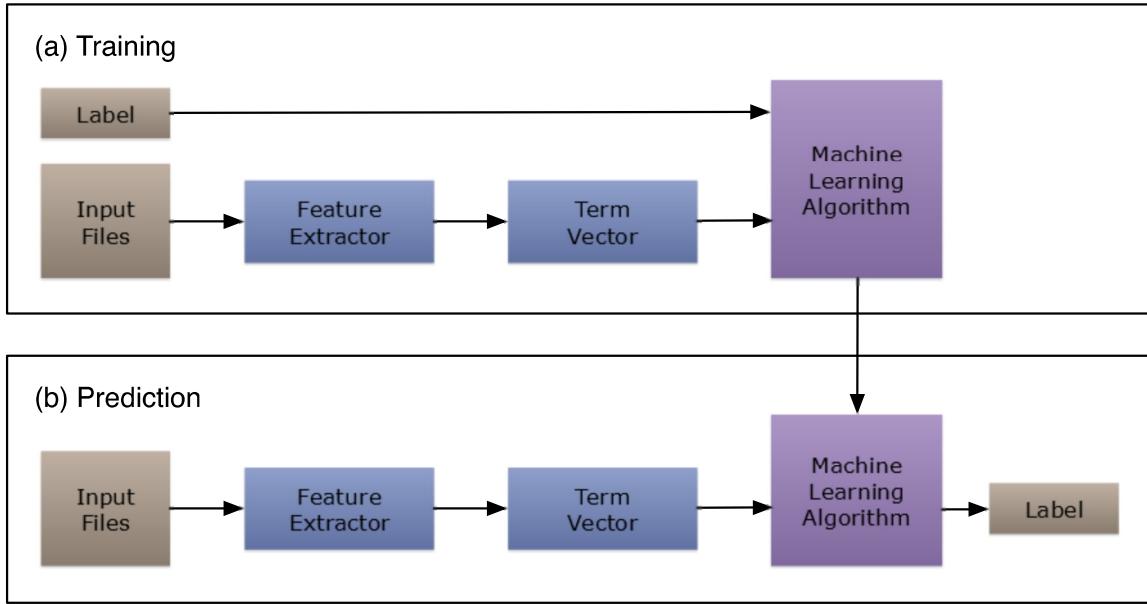


Figure 1: Supervised Text Classification

2.2 Sentiment Analysis

Sentiment analysis is a method of data mining to determine the attitude of a speaker or a writer with respect to some topic or the overall contextual polarity of a document. Many sentiment analysis techniques have been developed in recent years. The simplest classification is to classify text into either a positive or a negative sentiment category based on text classification. The basic approach is lexicon-based [17], which is to analyze tweets based on the words that the text contains. The texts are scanned and checked if some specific sentimental words are contained. It has been defined in a vocabulary that some words are positive and some are negative and each of them is assigned a sentiment score. The whole text will be determined based on the score. However, it is difficult to maintain a dictionary of key words to calculate the sentiment score. For this reason, some supervised and unsupervised algorithms are also developed

and used for text classification [16], such as Naive Bayes, Decision Tree, K-Nearest Neighbors, Maximum Entropy, and Support Vector Machines. For these machine-learning algorithms, in order to do the classification, sufficient labeled data needs to be fed into the classifier to train the classifier. Based on the training dataset, the classifier will build a probability model that is able to give a prediction of next input.

2.2.1 Algorithms

2.2.1.1 *Naive Bayes*

Naive Bayes is one of the most common supervised classification methods that can be used to perform text classification [18, 19]. Before that, we need first to have a look at what the feature vector is. In order to perform classification, we need to select features from the data first. For text classification, the feature vector is also called the term vector, which is the most important structure during the training and classification process. All tweet texts will be transformed to term vectors to be processed by classifier. Usually, term vector is generated based on a unique vocabulary, which is generated from the training dataset, and there are no duplicate words in the vocabulary. The size of the term vector is the size of the vocabulary. There are two types of Naive Bayes implementations: Naive Bayes – Bernoulli and Naive Bayes – Multinomial. The main difference between them is the way in which features are extracted from the documents. Let us take Bernoulli implementation as example, for a sentence like a tweet, a term vector will initialized with all elements equal to zero. Then check each word in the vocabulary to see if the word exists in the tweet. If it exists, then mark the corresponding element in the term vector to 1, if not, mark the corresponding element in the term vector

to 0. In that way, if the vocabulary is big enough, every single tweet can be represented using a term vector with 0s and 1s.

Generally, in text classification, we will ignore the order of words in the document. Instead, we only consider the presence or absence of the single word, for instance, whether a word in the vocabulary is included in the document or not. This model is called a bag of words. It is like we throw all words in a bag and they could be in any order in the bag.

The element of term vector does not only represent the presence or absence of a word, it can also represent the frequency of the word.

We can think every term vector as an n-dimension point in an n-dimension coordinate system, where n is the size of the vocabulary. For the training data, we can treat it as several classes of a bunch of n-dimension points. Then the text classification problem becomes traditional points classification problem, though the dimension might be very large.

For example, if we have a document D, and we have classes C, which contains some classes. Then in order to get which class that the document D belongs to, we just need to compute the posterior probability $P(C|D)$, and choose the largest one. $P(C|D)$ can be computed by Bayes' Theorem:

$$P(C|D) = \frac{P(D|C)P(C)}{P(D)} \propto P(D|C)P(C)$$

where the prior probability and likelihood can be computed from the labeled dataset.

Here is how these two algorithms work. Every class is equal probable, so we can easily get the prior probability $P(C_i)$. Let V be the vocabulary, and w_j is the j-th word in

V , so from the training data, we can get the probability of w_j belongs to class C_i , i.e. $P(w_j|C_i)$.

Then assume that S_i is the i -th tweet in my test dataset, and T_j is the term vector for this tweet. As we have known, T_j contains 0s and 1s that mean if the corresponding word in the vocabulary exists in tweet S_i . Then the probability of tweet S_i belongs to class C_i is

$$P(S_i|C_i) \propto P(T_j|C_i) = \prod_{j=1}^{|V|} [T_j P(w_j|C_i) + (1 - T_j)(1 - P(w_j|C_i))]$$

This equation is easy to understand, it means the multiplication of the probabilities that this tweet is composed by words in the vocabulary.

2.2.1.2 *Support Vector Machine*

The support vector machine (SVM) is also one of the well-known supervised machine learning algorithms to perform text classification [19, 20]. The main point of support vector machine is to find a linear separator in the search space that can best separate the different classes. As Figure 2 shows, there are two classes A and B. Three hyperplanes between them, I, II and III, separate them into two classes. We will choose the hyperplane which has the largest normal distance of any of the data points as the best separator, which is hyperplane I in the following figure.

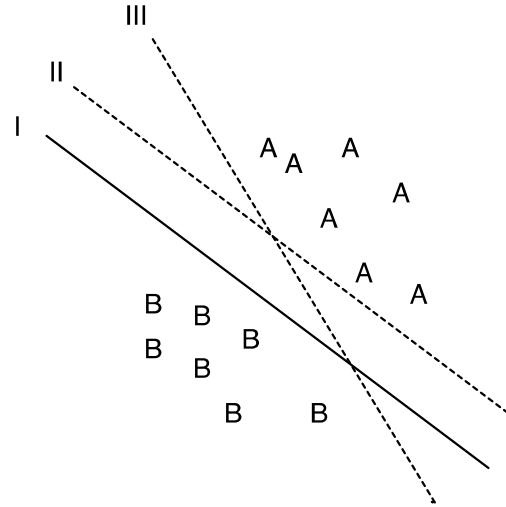


Figure 2: Support Vector Machine on Classification

2.2.1.3 *Decision Tree*

A decision tree [19, 21] is a simple flowchart that selects labels for input values. This flowchart consists of decision nodes, which check feature values, and leaf nodes, which assign labels. To choose the label for an input value, we begin at the flowchart's initial decision node, known as its root node. This node contains a condition that checks one of the input value's features, and selects a branch based on that feature's value. Following the branch that describes our input value, we arrive at a new decision node, with a new condition on the input value's features. We continue following the branch selected by each node's condition, until we arrive at a leaf node that provides a label for the input value.

For the text classification, the decision nodes could be the feature word we selected, and the leaf nodes could be the categories.

2.2.2 Twitter Sentiment Analysis and Stock Price

In the past, public opinions and public sentiment have been suggested and generally accepted as possible indicators of a company's stock behavior [24]. In recent years, there has been an emergence of number of social networks in which a person can share their emotions, feelings, status or locations, and a surge in their popularity. More specifically, Twitter has become one of the most important sources of public sentiment on various topics about companies, products, movies and many others. On Twitter, although each individual user post, which is called a tweet, is limited to only 140 characters, the millions of tweets could represent public opinions and public sentiment at a certain extent [25]. Among all their opinions, the opinions of specific companies have grown in prominence, and it can be argued that Twitter posts by a user are an accurate representation of a said user's opinions and thoughts.

It has been discussed that there is a correlation between public opinions and the stock price [24]. One of the stock price indicators is the market, and customers' behavior will have a significant impact on the market. So generally, the public sentiment about a company and its products is proportion to its stock price behavior. However, little work has been done to actually expand on this topic. There, for instance, is no time frame as to how long it takes for a company's stock to act in accordance to its associated twitter sentiment, or if there is any demographic on Twitter that does a slightly better job of predicting a stock's behavior than another.

2.3 Current Distributed Data Processing System

As we mentioned, the MapReduce is not the only solution for Big Data problem. There are lots of other large-scale data processing solutions and frameworks. Spark [6, 7], SCOPe [8], Dryad [9] are some other general-purpose systems for large-scale data processing. Besides the traditional MapReduce, there is a next generation MapReduce version called YARN in the latest version of Hadoop. In addition, Mesos [10] is a resource manager for multiple systems including MapReduce, Spark and MPI to share cluster resources. There are also some frameworks designed for specific computing. For instance, HaLoop [11] is a MapReduce framework specifically optimized and designed to perform iterative computing. Pregel [12] is famous for its large-scale graph computing. Storm [13] is special designed to perform real-time streaming data processing.

2.4 MapReduce Model

MapReduce is a parallel data processing paradigm, which was originally developed by Google Inc. The MapReduce model contains two phrases, the map phrase and the reduce phrase. Tuple, which contains a key and a value, is the key data structure in MapReduce. All data are transferred in tuple structure through the whole process of MapReduce task. The input to the program is organized in a $\langle k_1, v_1 \rangle$ pairs. The input will flow into the map function, and the output format of the map function is a list of $\langle k_2, v_2 \rangle$ pairs. These outputs will flow into the reduce function, while before that, they will first be organized by key k_2 . Each reduce function will only accept pairs with the same key, and a list of values of this key. After these pairs are processed, the output of

reduce function will be in the format of $\langle k_2, v_3 \rangle$ pairs. This process can be presented as following equations:

$$map(\langle k_1, v_1 \rangle) \rightarrow list(\langle k_2, v_2 \rangle)$$

$$reduce(\langle k_2, list(v_2) \rangle) \rightarrow \langle k_2, v_3 \rangle$$

In general, there are three stages of a MapReduce task: the splitting stage, the mapping stage, and the reduction stage. Here we will have an overview of the three stages in detail.

1. Splitting Stage

When an input dataset comes, the first step is to split it into multiple chunks. These chunks can be processed in parallel. From the view of this paper [14], the suggested size of a chunk is equal to the size of the L1 cache available to the hosting CPU core. Also, there is function called *splitter* that user can overwrite to define their own split rules. Usually, this step will be completed by Hadoop automatically, we do not need to configure a lot for this stage.

2. Mapping Stage

This stage will read the data chunk and partition the data into tuples. Usually, for a text input file, the input of map function will be every single line in this file. Through the user defined map function, each line will be processed individually and emitted in the format of key and value. The outputs of map functions are the partial results, and these partial results will be shuffled and sorted by key after they are emitted into the context. We can use our own defined *key_comparator* to define the rules of equality of two keys. Also,

these shuffled and sorted are called intermediate data and stored in local file system to be used later in the reduction stage.

3. Reduction Stage

All the aggregation happens in this stage. This stage begins after all map tasks have completed. Each reduce function will accept a list of values with the same key. The inputs are from intermediate data. For the list of values, we can define our own process, and finally, the final outputs of reduce function are emitted at last. The output is still in tuple format, where the key is the input key, and the value is usually the aggregation of all the values in the input. All the output tuples will be written into files in the distributed file system.

This model can be used to apply on large amount of data. The data will be organized into tuples and passed to map function and reduce function. Both the map function and reduce function can be run parallel on multiple machines throughout the cluster.

2.5 An Overview of Hadoop MapReduce Cluster

Hadoop is an open source Big Data framework with the implementation of MapReduce from Apache software foundation. It is written in Java and it provides us a high-level programming model for MapReduce. In this section, we will have an overview of the typical infrastructure of a Hadoop cluster, the Hadoop Distributed File System [15], and the MapReduce in Hadoop.

2.5.1 Hadoop Cluster Infrastructure

Figure 3 shows the typical two level network architecture for a Hadoop cluster. We can see the cluster is organized into racks. There are typically 30-40 nodes in a rack, and each rack is connected with a switch to a core switch or router.

In general, there are two kinds of nodes in a Hadoop cluster: the name node (also called the master node) and the data node (also called worker node). The name node manages the file system, which is known as Hadoop Distributed File System. The name node knows the data nodes on which all the blocks for a given file are located. The data node is the worker of the file system. When a task begins, the data node will find and retrieve the data from the file system and process the data. The file location is got from the name node. Also, the data node needs to report back to the name node periodically with the block locations they are storing.

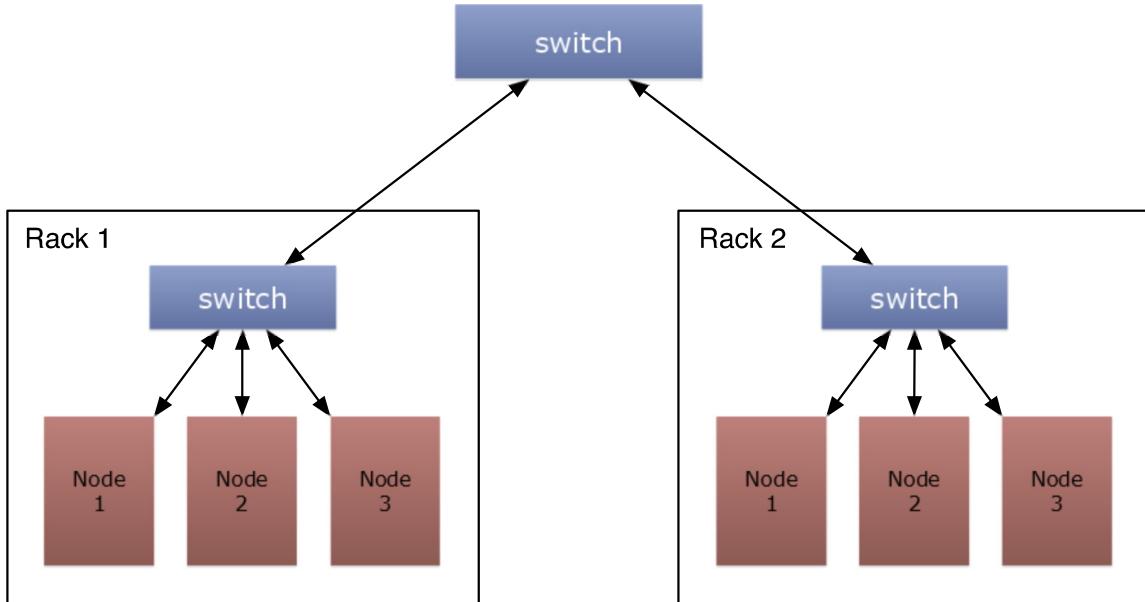


Figure 3: Hadoop Cluster Infrastructure

In general, a Hadoop cluster provides a whole runtime for MapReduce, plus a distributed file system over the whole cluster.

2.5.2 Hadoop Distributed File System (HDFS)

The Hadoop Distributed File System, also known as HDFS, is another important feature included in the Hadoop cluster. We have known from the previous part that the HDFS contains two kinds of nodes: the name node and the data node, and the name node is to manage the file system, while the data node is to store the data.

In a typical HDFS, the data could be stored on any one of all data nodes, which means HDFS will divide the data into fixed size blocks, and spread them across all data nodes in the cluster. Also, Hadoop is designed to run on commodity hardware, there could be file failures. So each block will be replicated three times, of which two replicas are placed in the same rack and the other one replica will be in the outside rack. The name node knows where the blocks are located and tells the location to the data node once needed.

2.5.3 MapReduce

Hadoop mainly contains two parts: the HDFS and the MapReduce. On the top of HDFS, the Hadoop MapReduce is the runtime and execution framework for MapReduce program. The typical Hadoop MapReduce contains two kinds of nodes: a single master node called JobTracker, and the worker node called TaskTracker. Usually the TaskTracker runs on the same nodes that HDFS data nodes run on.

When a task is assigned or submitted from client, the JobTracker is responsible for taking it in and assigning TaskTrackers with tasks to be performed. Here we have a terminology called Data Locality, which means the JobTracker will first try to assign tasks to the TaskTracker on the data node where the data is locally stored. Otherwise, the JobTracker will at least try to assign tasks to TaskTrackers within the same rack. In addition, the JobTracker is responsible for failure tolerance. If a TaskTracker fails, the JobTracker will detect that and re-assign the task to another TaskTracker. This ensures the stability of the whole cluster.

TaskTracker runs on each data node and is ready to receive tasks from JobTracker. In order to let the JobTracker know the running status, the TaskTracker will send heartbeat to the JobTracker indicating that it is alive. Also, the TaskTracker will send task running statistics to JobTracker periodically.

Figure 4 shows the relationship between the name node, the data node, the JobTracker, and the TaskTracker. In summary, a typical data execution flow in Hadoop cluster is as follows:

- (1) The program (also called the client) initializes a MapReduce job and submit it to the JobTracker;
- (2) The JobTracker will get the file location information from the name node, and try to put the client program in the HDFS. After that, it will try to assign tasks to TaskTracker based on data locality;
- (3) The TaskTracker receives the task and start the mapping phrase. The output of map function will be stored on local file system as intermediate files. After

that, the intermediate results will be passed to reduce function and the final output of reducing phase is written into HDFS.

(4) After the reducing phase, the job is completed.

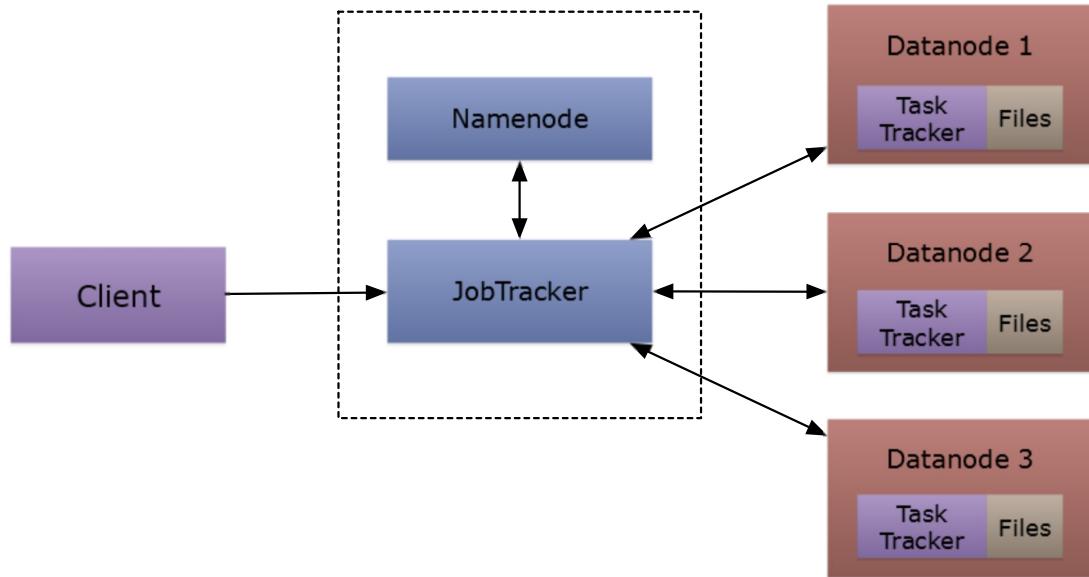


Figure 4: Typical Components of a Hadoop Cluster

3. NAIVE BAYES FOR TWITTER SENTIMENT ANALYSIS

In this chapter, we will have an overview of the prototype system to find the correlation between the geographical sentiment on Twitter and the stock price using Naive Bayes sentiment analysis algorithm, including the data source introduction, the system methods and implementation, and finally some experimental results.

3.1 Geographical Sentiment and Stock Price

Some researchers have shown that the public sentiment and stock price could be correlated [24]. The micro blogging site, Twitter, has been used to obtain the public's mood for elections [26] and some other particular events, and the idea of using sentiment found on Twitter and comparing it to a stock's behavior has been explored. It has been concluded that there is a strong indication that it may be possible to predict stocks prices to a certain degree using sentiment analysis. However, location is also an important aspect of social network, and the location's sentiment could be an indicator of stock price's behavior. Work has not been done on correlating a specific location's sentiment affecting a company's stock behavior.

Twitter doesn't provide many demographical options, such as gender and ethnicity. Research has been done to try to find the users' demographical information [24, 27] by using the user's self-reported geographic location and seeing if the demographic information can be accurately represented by Twitter; it was found out that Twitter sampling that represents 1% of the United States population is biased. No recent papers on geographical sampling expressed explicitly that sampling was biased [28]. Other work

has been done on geographic locations as well, this was a worldwide study and included geographical locations on user's offline twitter locations [27]. This section applies the geographical technique from [28] to obtain locations to locate the most populous cities.

3.2 The Source Data

It was found out that due to the privacy policy of Twitter, public tweets dataset cannot be obtained. In addition, all the public tweets dataset do not have geographical information included, which makes it impossible to use the existing datasets.

This system uses Twitter Streaming API to collect the real time data. The input streaming data is from Twitter's Streaming API, which is offered by Twitter to give developers low latency access to Twitter's global stream of Tweet data. We can get all the global public tweets by passing keywords. The data we receive is in JSON format. Every tweet is in JSON format including all the information about this tweet. Once connected to the Streaming API, the data will flow to the system constantly with long connection. The Twitter Streaming API accepts various filters, for instance, keywords, and languages. With the keyword filter, the system can only get the tweets about a specific company or product. The returned data may not have geographical information; in this system only the tweets with geographical information will be used. With the language filter, only English tweets are accepted.

The financial data, which is the stock price, is collected from Google Financial API [29], at 2 minutes intervals.

3.3 System Methods and Implementation

This section introduces the methodology and the implementation of the system.

The prototype system will be able to collect tweet data about a specific topic, process the data, and finally give the geographical sentiment data. Also, the system is designed with three layers, the data collection layer, the sentiment analysis layer, and the plotting layer.

3.3.1 Three Layer Design

The following figure shows the architecture of the three layer design.

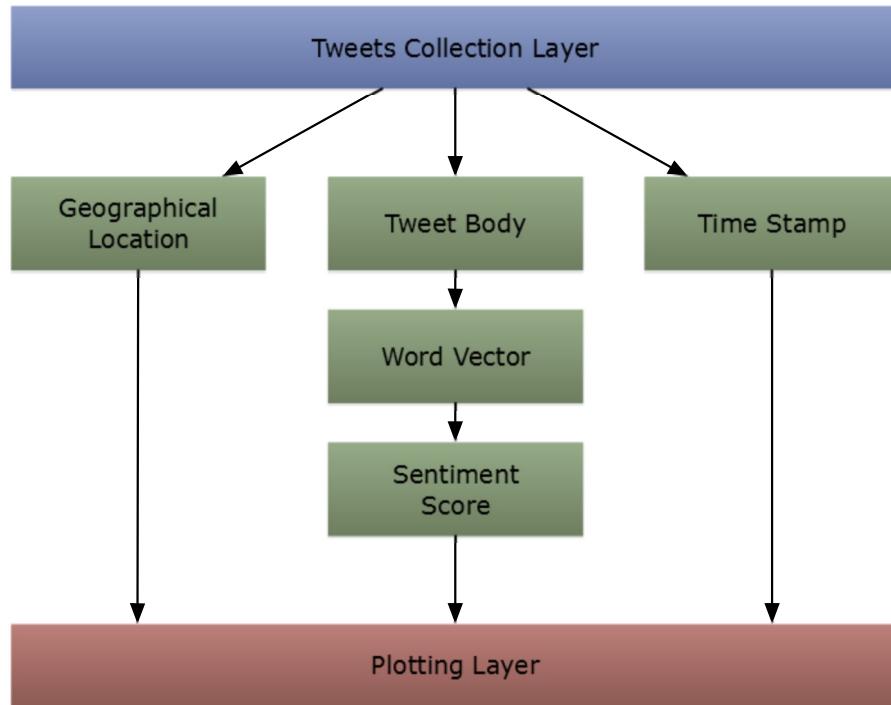


Figure 5: Three-Layer Architecture of the System

The Tweets Collection Layer is to collect the tweets returned by the search query about a keyword using Twitter Streaming API. The raw data contains a lot of information but the system only extracts the geographical location information, the tweets body, and

the time stamp. For the tweets body, it will be processed by Sentiment Analysis Layer, which will give a sentiment score based on Naive Bayes approach, and the result will be sorted by geographical information. At last, the Plotting Layer will plot graphs as needed for us to do analysis.

3.3.2 Stock Market Sector

A share of stock is literally a share in the ownership of a company. When a person buys a share of stock, that person is entitled to a small fraction of the assets and earnings of that company. This means that the more a company earns then that person will invest in stocks since they hope the company will do well. A source of company earning is from selling products and services; which is affected by consumer purchases. This means, however, that the more positive the consumer's view of a company and its potential earnings, there may be a higher demand for the stock of said company, and, in turn, the opportunity to raise the price of a company's stock is presented. Essentially, this means that a company's reputation among consumers can influence its earnings both indirectly and directly.

The first challenge of our experiment was choosing a company to study; this company's stock price and behavior would have had to be tracked and recorded, and later graphed in order for us to analyze it alongside the graph of its associated Twitter sentiment.

As the above requirements are somewhat specific, and as there are thousands of companies trading stock right now, we had to narrow down the search in order to get more suitable candidates from Standard & Poor's 500. A specific sector was chosen, and

companies falling into that sector were considered. Our sector of choice was food and restaurants, simply because there was a limited possibility of advertising-related post and non-relevant posts to be streamed into our dataset, unlike technology companies or TV companies that have more of an online presence. From previous works, it was very common to use the Standard & Poor's 500 to compare stock price and mood or sentiment, so we decided to draw our companies from that list when deciding which one we would actually choose.

We selected tweets based on company name keywords, stock ticker symbol, and uniquely named popular food items for example Frappuccino or Happy Meal. Stock ticker symbol have been commonly used to search stock related information about a tweet. We decided to take this into consideration since previous studies have used this technique. Our goal is to find as many stock related tweets as possible in order to represent the full sentiment of a company on Twitter.

3.3.3 Geographical Location

As mentioned earlier, there has been some work linking tweets and the stock market already. It is somewhat accepted in the research community that Twitter, among other social media sites, can, in some situations, serve as a replacement for the general public, and polls and questionnaires can now be performed with online media as a platform. However, because of the relatively new emergence of social media, there is still much growth for improvement in terms of research, data mining, and a number of other things.

One of our main objectives is finding out whether tweets from one location have a better correlation to a stock's price than another, and possibly investigate why. Our experiment relied on not only taking the general sentiment found on Twitter, but also dividing this sentiment into different geographical locations. Twitter users have a choice to report their geographical location on their profile. Locations were collected when streaming the tweets, and kept a tally on how many tweets each location had. Once all the streamed tweets were processed, and all of the self-reported locations had a count of the tweets from users of that location, we simply chose the five states that had the highest number of tweets. Originally, it was proposed to only filter out tweets from five pre-specified cities – to be more specific, the five cities that contribute most to the American GDP – but previous work suggests that Twitter population is not necessarily representative of the American population, and any of the cities we proposed may not have had as large of a Twitter presence as we originally thought. In addition, we also found that if we set the location to a specific city, the amount of tweets we can get is too small to do some reasonable analysis. Therefore in this system, we expanded the location box to state.

3.4 Experimental Results

The system has collected all the tweets about Starbucks and McDonald's for four days from 9:30AM to 4:00PM EST. In this section, we take Starbucks as an example. Before that, we will first have an overview of the result representations.

3.4.1 Result Representations

The first definition is bullishness score (B_t) for a time period t from paper [30] given as:

$$B_t = \ln \left(\frac{1 + N_{positive}}{1 + N_{negative}} \right)$$

In this equation, for each time interval t, the total number of positive tweets is aggregated as $N_{positive}$, and the total number of negative tweets is aggregated as $N_{negative}$. In this system, this score will be used to plot the graphs.

The expected output is geographical sentiment (bullishness) scores over time about a specific company or product. The result is sorted by states. The goal of our system is to track the changes of both geographical sentiment scores and stock prices so as to find if there is a strong correlation of them. The best way for us to understand the result would be plot them into some.

The NASDAQ Stock Market operates from 9:30AM to 4:00PM EST. So all the stock price data and tweets are collected between 9:30AM and 4:00PM EST. This time period is divided by 10 minutes interval and for each interval in the sentiment graphs, the bullishness sentiment score is used.

In addition to the plotting, this system also calculates the correlation coefficient in order to have a measurement of how similar the two graphs are. The correlation coefficients are calculated based on the following equation:

$$CORREL(X, Y) = \frac{\sum(x - \bar{x})(y - \bar{y})}{\sqrt{\sum(x - \bar{x})^2(y - \bar{y})^2}}$$

In this equation, \bar{x} and \bar{y} are the sample means of points array X and Y . For the result, a correlation coefficient of +1 indicates a perfect positive correlation, while a correlation coefficient of -1 indicates a perfect negative correlation.

3.4.2 Analysis of the Results

The figure below is the preliminary results for a single day about Starbucks. In the graph, each of the top five states are above, in the order of California, Illinois, New York, Texas, and Washington. Each of the above graphs has a timeline of about six and half hours, divided by 10 minutes interval, and an average sentiment (bullishness) score of that time period. The graph in the shadow is the stock price line graph from Google Finance, which is a free service for us to get the historical stock price data and graphs.

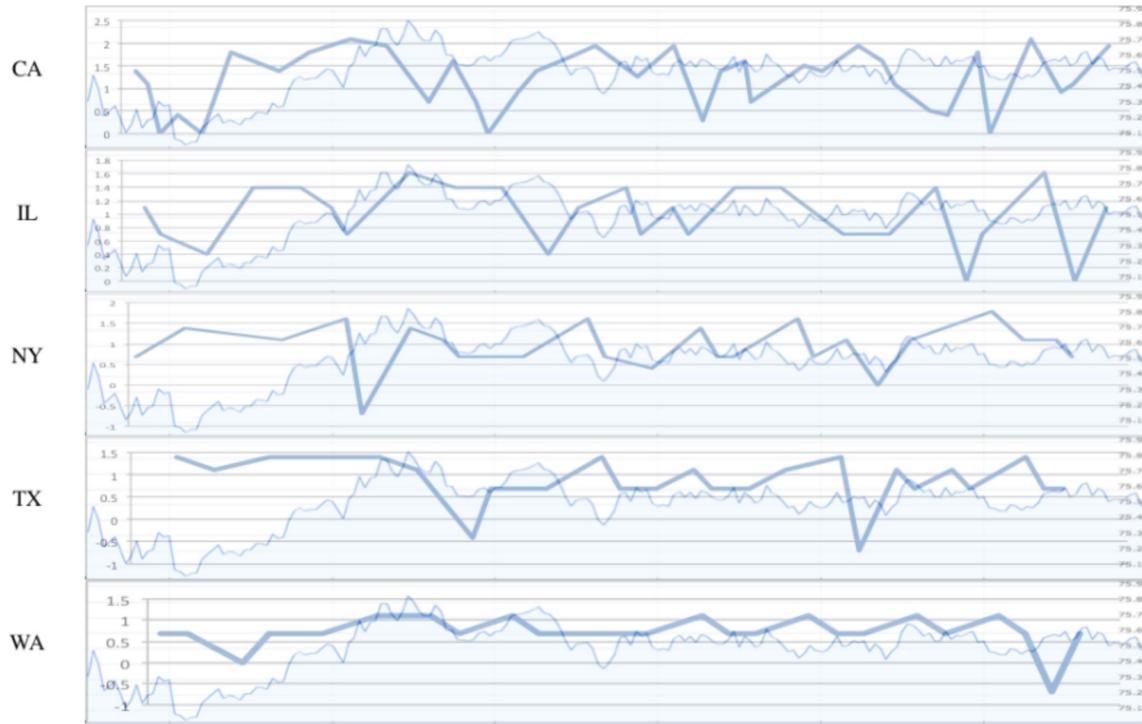


Figure 6: Sentiment of Starbucks Corporation by City

This result is just for a single day. It seems that there is no strong correlation between the geographical sentiment and the stock price. As a comparison, the following figure shows the other three days' data.

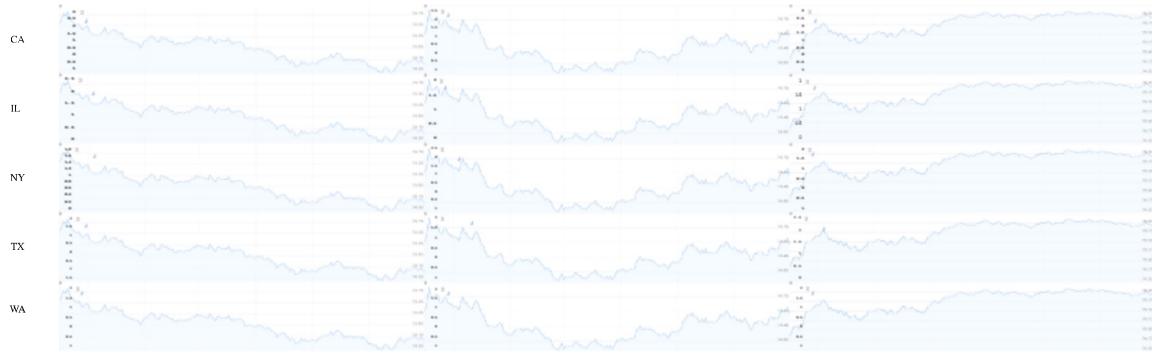


Figure 7: Three Days of Sentiment of Starbucks Corporation by City

In order to have a clear view of their correlations, the following table shows the correlations coefficients.

Day	1	2	3	4	AVG
CA	0.179	0.138	-0.033	0.127	0.103
IL	0.031	0.442	0.693	-0.027	0.285
NY	-0.110	-0.121	0.129	-0.021	-0.031
WA	0.337	-0.184	0.224	-0.365	0.003
TX	-0.281	-0.131	0.017	-0.199	-0.149

From the graphs in Figure 6 and Figure 7, it hardly shows correlations between the shadow lines and the single lines. Sometimes, the two lines may have the same trends, such as the Washington State in Figure 6, and the first half part of the Illinois State in the first graph in Figure 7. However this is not always true, the Washington State has a bad match in the third graph in Figure 7 and the Illinois State also has a bad match in the

second graph in Figure 7. In order to know the percentage, or the similarity of the match, we calculated the correlation coefficients.

From the correlation coefficients table, the New York and the Texas States have 75% negative correlation in all four days, while the other three states in general have a positive correlation. Also, the Illinois State has the strongest correlation, while the Texas State has the weakest correlation. This conclusion is about the same as we observed from the graphs.

Based on what we observed and calculated, we can say that for Starbucks Corporation, the Illinois State has a strong correlation between the Twitter sentiment and the stock price. In addition, although the headquarter of the Starbucks is located at Washington State, the result does not show that the Washington State has a very strong correlation. In contrast, the Washington State has a very weak but not negative correlation. In comparison, the Texas State does not have such a strong correlation.

4. NAIVE BAYES MAPREDUCE IMPLEMENTATION

In this chapter, we will introduce the methodologies used and Naive Bayes MapReduce implementation in this thesis. First, we have an overview of the data processing methods we used to pre-process input data before it is applied to training stage. After that, we introduce the parallel implementation of Naive Bayes based sentiment analysis algorithms in MapReduce. In addition, we review the evaluation metrics and scenarios used for this thesis.

4.1 Methods Overview

Before we enter into each part, let us have an overview of the whole method, as the following Figure 8 shows.

As the graph shows, the input data is first pre-processed to remove some noises that are not useful for sentiment analysis. We implemented a MapReduce system to perform sentiment analysis based on Naive Bayes classification algorithm in parallel. The pre-processed data will be fed into the system and ran under different metrics, including the size of input data, number of nodes, use of global counters, and hardware configurations. Based on the results, we combine the analysis with some MapReduce performance factors and try to build the MapReduce performance prediction model.

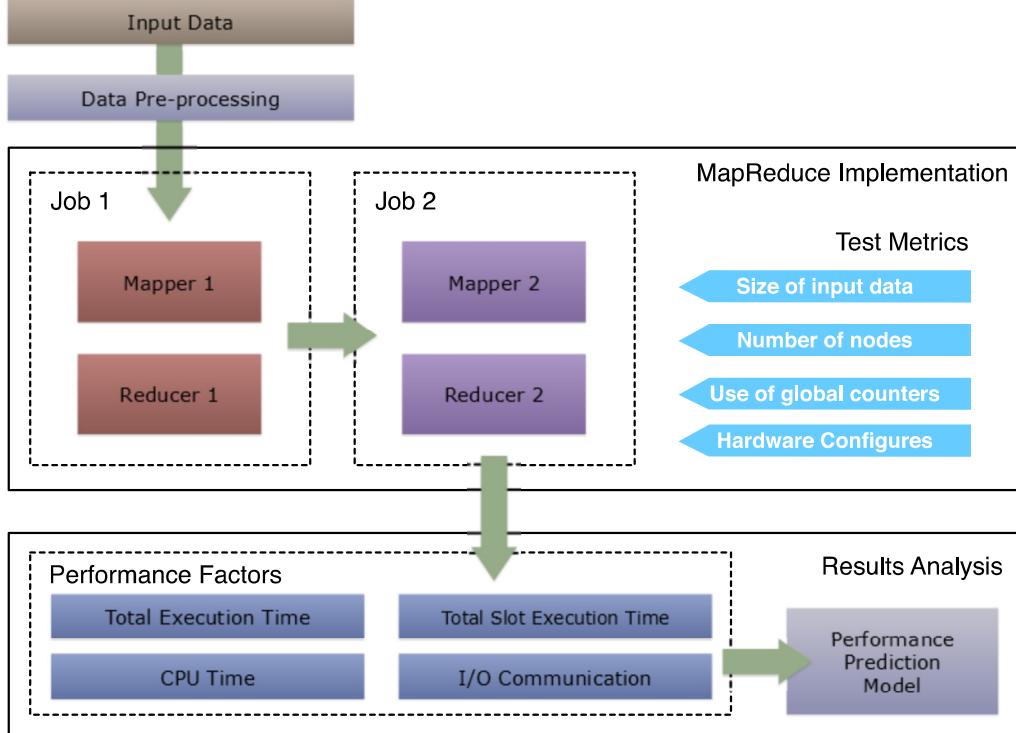


Figure 8: Whole System Method Overview

4.2 Data Pre-processing

We chose Twitter as the source of our data because the Twitter is one of the most popular social networks in the world. According to Twitter's IPO filling [22], it has more than 200 million monthly active users who tweet 500 million tweets per day in 2013. Another important reason is tweets are almost all about text, while on Facebook or other social networks the media could be videos, pictures or sharing web links. Twitter allows people to publish texts with 140 characters limitation. When talking about something, people intended to talk more specifically and with more feelings or attitudes.

However, Twitter also has some drawbacks. As we know, tweet is not formal text. Besides the limitation of 140 characters, people also like to use a lot of abbreviations,

slangs, emoticons etc. to try to use the least words to express as much feelings as they can. Also, Twitter allows us to post URL links in the text, to use ‘@’ to refer somebody, and to use ‘#’ to give tags. This information is not very useful for the sentiment classification. So the following preprocess is performed to each tweet at the very beginning.

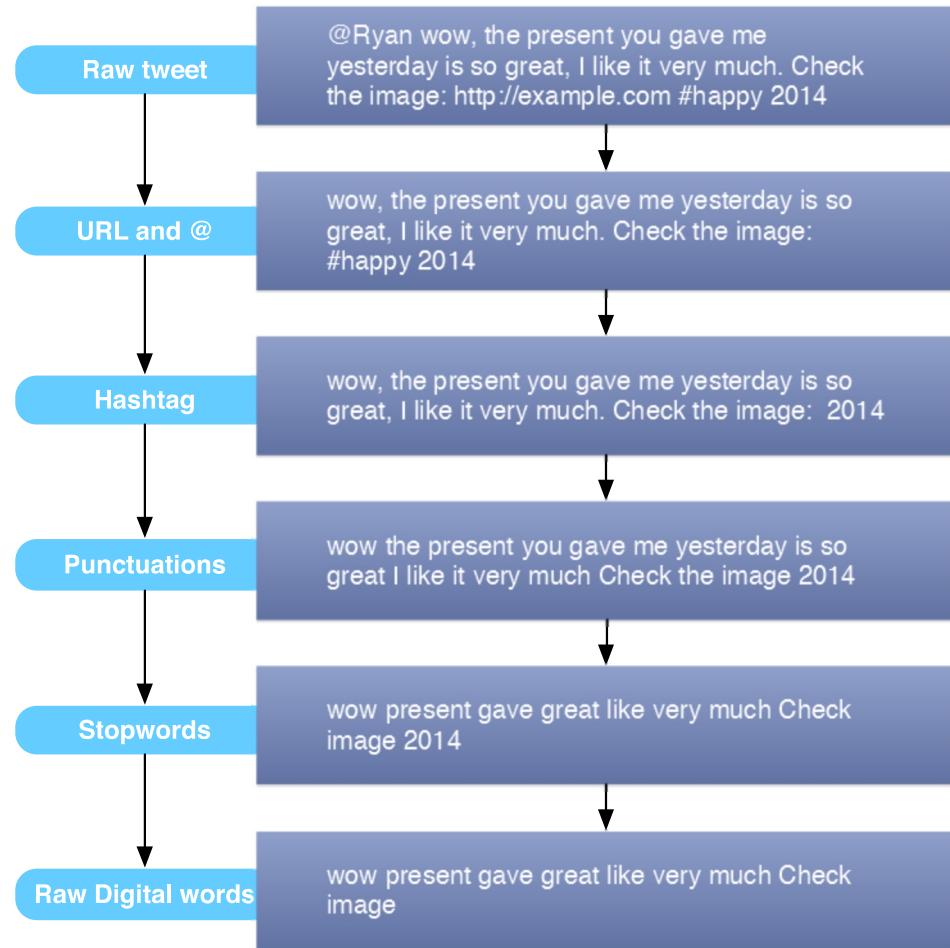


Figure 9: Tweets Data Pre-processing

1. URLs and ‘@’ Removal

The first step is to remove URLs and the word starts with ‘@’ symbol. We will not track the content of the web links, so the URLs are deleted. The ‘@’ symbol always

has a username followed, which is useless, so the entire word starts with ‘@’ could be removed.

2. Hashtag Removal

The word starts with ‘#’ is hashtag. Hashtag is different from other words; it gives a tag, or a topic to the tweet. Usually, the tag is talking about the topic people saying about in this tweet, not about people’s attitudes. This word might provide some information but not that important. So I decided not to remove the entire word, but to just remove the ‘#’ symbol, and treat the tag as a normal word in tweet.

3. Punctuations Removal

We do not need punctuations as the features, they are just symbols to separate sentences and words, so they should be removed.

4. Stopwords Removal

There is a kind of word called stopword [24]. They are some common function words in a sentence, like ‘a’, ‘the’, ‘and’, ‘to’, ‘at’, etc. These words seem to be useless for sentiment analysis, and also these words appears a lot of time in English, if we use term frequency to determine if a word is informative, these words will account for a large proportion. So these words should be removed.

5. Digital words Removal

Some words start with a digit, like ‘1990’, ‘4:00pm’ etc. These words also have no relationship with attitudes or feelings. So these words should be removed.

After the pre-process of dataset, all tweets will only have some plain words. Through this pre-process, noises are removed so that we can build better vocabulary and have smaller dimension of the term vector.

4.3 Parallel Naive Bayes in MapReduce

We have already introduced the sequential Naive Bayes algorithms for text classification based sentiment analysis algorithms in chapter 2. MapReduce is another new parallel programming model so that we need to adjust these algorithms to fit into MapReduce model. In this section, we will introduce the parallel implementations of Naive Bayes based sentiment analysis algorithm in MapReduce model.

4.3.1 Algorithm Implementation

Let us first have a look at the equation of Naive Bayes for sentiment classification.

$$P(C|D) = \frac{P(D|C)P(C)}{P(D)} \propto P(D|C)P(C)$$

We can see from the equation that the most two important values we need to compute the posterior probabilities are class prior probabilities $P(C)$ and likelihoods $P(D|C)$. For a training set, it seems that in order to get these two values, we need to have some counts. Here we use C to refer to the label of the tweet, and X to refer the features (words). What we needs are:

1. Number of tweets: the number of all tweets.
2. Number of C : the number of unique labels.
3. Number of $C = c$: the number of tweets with label c . This should be a global variable and can be used with the previous number we have to compute label prior probabilities.
4. Number of X in $C = c$: the number of words that are in all the tweets with label c .

5. Number of $X = x$ in $C = c$: the number of word x that is in all the tweets with label c . This number and the previous number we have are combined to compute the probabilities of a single word x under a label c .
6. $|V|$: the total number of unique words in all tweets. This is actually the size of our vocabulary.

Using the above data, we can get the prior probabilities $P(C)$ and likelihoods $P(D|C)$.

It is obvious that only one mapper and reducer job is not enough. And our goal is to build the prediction model, which means we only need to test the training phase. In general there are two jobs: the word job and the label job. For each job there is a map function and a reduce function. The word job is to compute the likelihood of each word, and the label job is to compute the prior probability of each label. Let us take a look at each job in details.

1. The Word Job

This job is mainly about word counting. The reason we call it Word job is because this job will output all the statistics about every single word. As we know, the job will have a map function and a reduce function, and each function takes tuples. Here are the inputs and outputs of these two functions.

- **WordMapper**

Input: each line (tweet) of the text

Output: $\langle word, label \rangle$

Algorithm:

```

map(docID, line):
    label = line[0]
    text = pre_processing(line[1:])
    textArray = text.split(' ')
    for word in textArray:
        emit(word, label)

```

The mapper will receive every single line of tweet in the input dataset.

The line contains two parts, the first part is the label, and the second part is the tweet body. For the tweet body, we invoke a function called *pre_processing()* to perform pre-processing on the tweet as we mentioned earlier in this thesis before. We use two variables to store these two values: *label* and *text*. Then the pre-processed tweet will be split into every single word delimitated by space and stored in the variable *textArray*. Then we put every word and its label into a tuple and emit it to the context. For example, we have a tweet with label “positive”:

“*positive I like the weather today :)*”

After pre-processing we can get:

“*positive like weather today*”

Then the output of the mapper will be three tuples:

“*<like, positive>*”

“*<weather, positive>*”

“*<today, positive>*”

These outputs will be passed to the WordReducer.

- **WordReducer**

Input: $\langle word, label \rangle$

Output: $\langle word, label_1:count \; label_2:count \dots \; label_n:count \rangle$

Algorithm:

```
reduce(word, label[]):  
    |V|++  
    HashMap labelCounts;  
    for label in label[]:  
        if labelCounts has key label:  
            labelCounts[label]++  
        else:  
            labelCounts[label] = 1  
        for key, value in labelCounts:  
            output += "key:value"  
    emit(word, output)
```

During the shuffle and sort process, the output tuples from the WordMapper will be sorted by the key *word*. All the *labels* with the same key *word* will be passed into one reducer as an array. For example, we have these three values from the mapper:

$\langle like, negative \rangle$

$\langle like, positive \rangle$

<like, positive>

Then they will be re-organized as one tuple:

<like, [negative, positive, positive]>

In the reducer, because all the tuples with the same key *word* will be passed only to one reducer, so the first step is to add one to a global variable V , which is the size of vocabulary. This variable will have the count of all unique word in the corpus at last. Then for the label array, we use a HashMap to count every label. The HashMap will be defined as follows:

HashMap<String, Integer>

The key of the map is the label string, and the value with key is an integer number storing the count of the label. If the label is already in the HashMap, the value will be added by one, otherwise a new key associated with the label will be created.

At last, the key and value pairs will be put into an output string and emitted by the reducer with the *word*. For example, the output of the previous input will be:

<like, negative:1 positive 2>

Until now, the whole Word Job is finished. The output of the reducer will be written into a temporary file in HDFS and this file will be used later to get all statistics about training model.

2. The Label Job

This job is mainly about label counts. The reason we call it Label Job is because this job will generate some label statistics. Let us take a look at the mapper and reducer function.

- **LabelMapper**

Input: each line (tweet) of the text

Output: $\langle \text{label}, \text{word_count} \rangle$

Algorithm:

```
map(docID, line):
    label = line[0]
    text = pre_processing(line[1:])
    textArray = text.split(' ')
    emit(label, textArray.length)
```

The mapper will receive every single tweet as input. After pre-processing of the tweet, and split by space, we can get an array of the words in that tweet. In this mapper, we simply emit the word number along with the label. For example, if the input tweet is this:

“positive I like the weather today :)”

After pre-processing we can get:

“positive like weather today”

Then the output of the mapper will be a tuple:

$\langle \text{positive}, 3 \rangle$

This tuple means, this tweet is positive and it has 3 valid words or features.

The output will be written into intermediate files and passed into the LabelReducer.

- **LabelReducer**

Input: $\langle \text{label}, \text{word_count} \rangle$

Output: $\langle \text{label}, \text{tweet_number:word_number} \rangle$

Algorithm:

```
reduce(label, int[]):
    UniqueLabelNumber++
    for count in int[]:
        docNumWithLabelY++
        wordNumWithLabelY += count
    emit(label, docNumWithLabelY:wordNumWithLabelY)
```

The reducer will take all the word counts associated with the same label as input. The input is also re-organized by shuffle and sort process before passed into the reducer. In the reducer, there is another global variable named *UniqueLabelNumber*, which is the number of all unique labels, because each reducer will only accept one kind of label. Then for the word count array, we will iterate over each element and add the value to a variable *wordNumWithLabelY*, which is the number of all words in the tweets with label *Y* (the label *Y* here is just an example). At the same

time, another variable `docNumWithLabelY` will be added by one to calculate the number of tweet with label `Y`.

At last, the variable `docNumWithLabelY` and `wordNumWithLabelY` will be combined together and emitted out along with the label. For example, if the output is:

`<positive, 12:514>`

This tuple means there are 12 tweets that are belonged to positive category and of all that tweets there are 514 words totally. These results will be written into HDFS and used later.

Until now, let us have a look at what data we can generate from these two jobs. In this example, we take word “happy” as a feature word.

- **(*happy, positive:3 negative:1*)**

3 positive tweets having word ‘happy’, 1 negative tweet having word ‘happy’.

- **(*positive, 6:20*)**

6 positive tweets with 20 words totally.

- **(*negative, 5:16*)**

5 negative tweets with 16 words totally.

- **|V| = 25**

The vocabulary size, 25 unique words in all the tweets.

From the data above, let us see what we can get.

Prior probabilities:

$$P(\text{positive}) = \frac{6}{5 + 6}$$

$$P(\text{negative}) = \frac{5}{5 + 6}$$

Likelihoods (with Laplacian Smoothing):

$$P(\text{happy}|\text{positive}) = \frac{3 + 1}{6 + 25} = \frac{4}{31}$$

$$P(\text{happy}|\text{negative}) = \frac{1 + 1}{5 + 25} = \frac{1}{15}$$

After we get the prior probabilities and likelihoods, if we have a new tweet to be classified, we can simply apply the equation:

$$P(C|D) = \frac{P(D|C)P(C)}{P(D)} \propto P(D|C)P(C)$$

to get the posterior probabilities:

$$P(\text{positive}|\text{new_tweet}) \propto P(\text{positive}) \prod P(\text{word}_n|\text{positive})$$

$$P(\text{negative}|\text{new_tweet}) \propto P(\text{negative}) \prod P(\text{word}_n|\text{negative})$$

Then we select the category with a higher prediction rate as the result of our classification.

4.3.2 Algorithm Time Complexity

The time complexity of Naive Bayes based sentiment analysis algorithm is mainly determined by two processes: the computation of prior probabilities of classes, and the computation of likelihoods of feature each word $P(w_t|C_j)$. Let us assume that there are c classes, and the number of feature words, which is the size of vocabulary, is $|V|$. Then we can get:

- (1) Computing the prior probabilities of each class: $O(c)$

(2) Computing the term frequency of each feature word: $O(|V|)$

(3) Counting the number of documents/tweets: $O(1)$

(4) Computing the likelihoods of each feature word: $O(c \cdot |V|)$

In total, the time complexity of Naive Bayes based classification MapReduce algorithm on a single node is:

$$T(n) = O(c) + O(|V|) + O(1) + O(c \cdot |V|) = O(c \cdot |V|)$$

Assume that we have n nodes, and then the computation will be distributed to n nodes. So the time complexity of the parallel MapReduce Naive Bayes sentiment analysis algorithm on n nodes cluster will be:

$$T(n) = \frac{O(c \cdot |V|)}{n}$$

From the equation, we can see that the time complexity will be decreased while the number of nodes increases.

5. EXPERIMENTAL RESULTS

In this chapter, we will introduce some metrics we used in this thesis to test the performance, and the testing scenarios including our datasets and hardware configuration. We will analyze the testing results regarding to the relationship between the execution time and number of nodes, the input data size, the use of global counters, and the hardware configuration.

5.1 Metrics Summary

In order to have an overview of the execution status of a MapReduce job, we need to define some metrics to observe during the job execution. In this thesis, what we call the performance is actually defined by the time of execution. The shorter the execution time is, the better the performance is. There are a lot of other aspects that can affect the execution time, such as the number of nodes, the size of the input dataset, or the use of global counters in a MapReduce job, etc. So we need to change these metrics during the execution of the job and observe the changes of the results.

5.1.1 Terminologies

For a MapReduce job, we will first define some terminologies we used in this thesis.

(1) Total Execution Time

This is the time consumed during the whole job, which means it is counted from the point of job starting and to the point of job ending, including everything happened in the job.

(2) Total Slot Execution Time

This is a little different from the total execution time because this slot execution time is only about the time consumed by the mapper and reducer execution. This time is usually shorter than the total execution time, the difference is mainly about I/O transferring time, the network communication, and job scheduling time, etc.

(3) CPU Time

This is a little harder to understand, because this time will not change if a job does not change. This is the time consumed by CPU, which is used for processing instructions of a computer program. It is only related to the number of instructions in a program, which means if we parallelize a program, the total execution time might be the third of the original, but the CPU time would not change because the amount of instructions does not change but are executed in parallel.

(4) Global Counters

Hadoop framework does not have global variable sharing mechanism, while it provides a global counter mechanism. From the name we can know that it is mainly for counting. We may need to keep track of some counters like how often a certain event has occurred during the job execution. Because of this feature of global counter, we can use it to share some global values during the whole job, such as how many documents or tweets in total, how many unique words in total. Global counters are reported back to the job tracker from the task tracker repeatedly like heart beating.

5.1.2 Impact of Node Number

For a cluster, the number of nodes is one of the very important aspects that affect the whole performance of execution. As we introduced in the time complexity section, the time complexity and the number of nodes is inversely proportional. However we do not know the exact relationship between them. Therefore, the number of nodes is one of many metrics we will evaluate in this thesis.

5.1.3 Impact of Data Size

The data is the key to the program, and it directly affects the execution time because the program is to process data. It is obvious that in a sequential algorithm, the execution time should be linear to the size of data. However, in Hadoop cluster, the data is distributed and stored in Hadoop Distributed File System. There might be many other aspects that affect the relationship between the data size and the execution time. It is a cluster with many nodes, data may need to be transferred or re-organized at anytime. Plus the locality feature of Hadoop framework, the impact of data size still needs to be observed and analyzed.

5.1.4 Impact of Global Counters

In common programming practice, the global variables are usually avoided to use because they could bring a lot of sharing or synchronization problems. The global variable needs to be synchronized to avoid mutual operations. Also, the global variable will be allocated in shared memory, and any inappropriate operations would cause unknown problems. Usually, it is impossible to share global variables between nodes in a large cluster, or it will cost a lot of I/O and affect the performance of the program

seriously. Although Hadoop provides us a new mechanism to count some events across the whole cluster, we still need to know how costly it is during a MapReduce job execution.

5.2 Evaluation Scenarios

In this section, we will have an overview of the format of the dataset used in this thesis, and the hardware configuration, which is the Hadoop environment.

5.2.1 Dataset

The dataset we used in this thesis is in CSV format. In the data file, each line is a single input including two parts: the label and the body of the tweet. There are two categories in our dataset, positive and negative. We used number 0 to represent negative and number 4 to represent positive. Here is a snippet of the dataset:

```
0, I forgot my phone in my car but I'm too scared to go outside and get it.  
0, feels like crying, that's how sick I feel!!  
4, The SUN is shining!! I'm happy :)  
4, @hype6477 Good to hear all's cool. I've just sent you a DM hope it'll start to help you
```

As we can see from the snippet, there are two kinds of tweets: 0 for negative and 4 for positive. The tweet body in the second column is the original one with abbreviations, emoticons, or '@' symbol etc. in it. The size of dataset we used in this thesis is 0.25GB, 0.5GB, 1GB, 2GB, and 4GB.

5.2.2 Hardware Configuration

We used Amazon Elastic MapReduce (EMR) [23] as the experiment environment. As the official states, Amazon Elastic MapReduce is a web service that

makes it easy to quickly and cost-effectively process vast amounts of data. We can create a Hadoop cluster with different node types and applications (Hive, HBase, etc.) on the cloud. We can create any number of nodes in the cluster, and after that, there is SSH access to the master node. Using the SSH connection, we can upload our jar files, data files, execute Hadoop programs, or install our own dependencies on the master node. After the job is done, the EMR cluster can be terminated so that the computation resources can be assigned to others.

Besides the Amazon EMR service, there is another service called Amazon Simple Storage Service (S3) to provide secure, durable, highly scalable object storage. The advantage of Amazon S3 is that it can be used along with other Amazon Web Services, such as Amazon EMR. Although we can use SSH to connect to the master node and upload our dataset, this is not a very efficient way. If the dataset is very large, it will take a lot of time to be uploaded and we still need to put it into HDFS after it is uploaded to the local file system on the master node. Once the cluster is terminated, the data we uploaded will be erased permanently. Using S3 can solve all these inconveniences. S3 can be used directly in Amazon EMR as HDFS, which means, we can upload data files to S3 and we can get an URL to this file. This URL can be treated as an HDFS URL and shared by multiple nodes. In this case, the data will not be lost and the size of the data could be very large. The basic infrastructure of Amazon EMR is shown in the following figure.

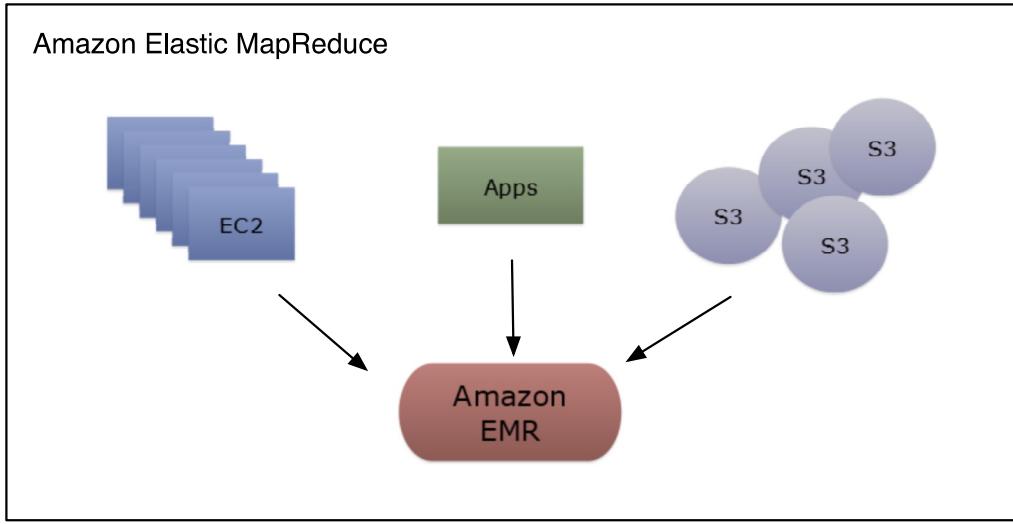


Figure 10: Basic Amazon EMR Infrastructure

In this thesis, the Hadoop cluster was created on Amazon EMR, and the data was uploaded to S3. There are several different kinds of nodes provided by Amazon EMR, and we used the M1 large type, with the following hardware configuration:

Model	vCPU	Memory (GiB)	Storage (GB)
M1 large	2	7.5	2 x 420

As comparison, we also used the M2 4xlarge type, with the following hardware configuration:

Model	vCPU	Memory (GiB)	Storage (GB)
M2 4xlarge	8	68.4	4 x 420

All these features and configurations make it a perfect choice for a Hadoop cluster. As for the number of nodes, we created several clusters with 1, 2, 4, 6 and 8 nodes, respectively to run different sizes of data and collect the results.

In the following part, we will demonstrate the running results of parallel MapReduce Naive Bayes based sentiment analysis algorithm on different sizes of data input with different clusters and configurations. Each metric will have visualized graphs and corresponding analysis.

To classify, the data sizes are 0.25GB, 0.5GB, 1GB, 2GB, and 4GB respectively and the numbers of nodes are 1, 2, 4, 6 and 8 respectively.

5.3 CPU Time

In this experiment, we used M1 large type of node. As we introduced before, the CPU time is the time consumed by instructions. For a specific job, if the size of data to be processed and the program itself are constant, the number of instructions that will be executed in CPU is constant. Therefore, for each input size, the CPU execution time should be constant no matter how many nodes are used for the job. The results are plotted in Figure 11 below.

In Figure 11, the vertical axis is the CPU execution time and the unit is minute. The horizontal axis is the number of nodes. The legend is the size of input data. There are five series in this figure, and each color represents a size of the input data. As we can see from the figure, if the input size is constant, the CPU execution time is constant as well, no matter how many nodes are used. This result verifies our hypothesis for CPU execution time.

From the result above, we conclude that the instructions executed in a job will not change. If the total time or total slot time changes later, then they must be affected by other aspects. This conclusion is the foundation of later hypothesis and analysis.

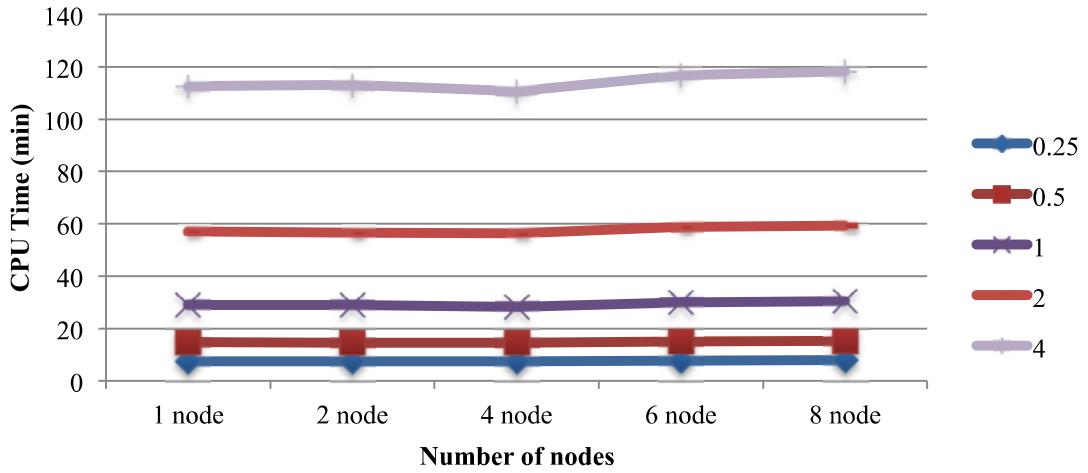


Figure 11: CPU Time and Number of Nodes

There is another important hypothesis that the amount of instructions executed by CPU should be proportional to the input data size because the main computing is about data processing. The following Figure 12 shows the relationship between the CPU time and the size of input data.

In Figure 12, the vertical axis is the CPU execution time and the unit is minute. The horizontal axis is the size the input data. The legend is the number of nodes. We can see that the five lines overlap because the CPU time is not affected by the number of nodes. And what the most important is the CPU time and the size of input data are linearly related, which verifies the hypothesis we made before is correct.

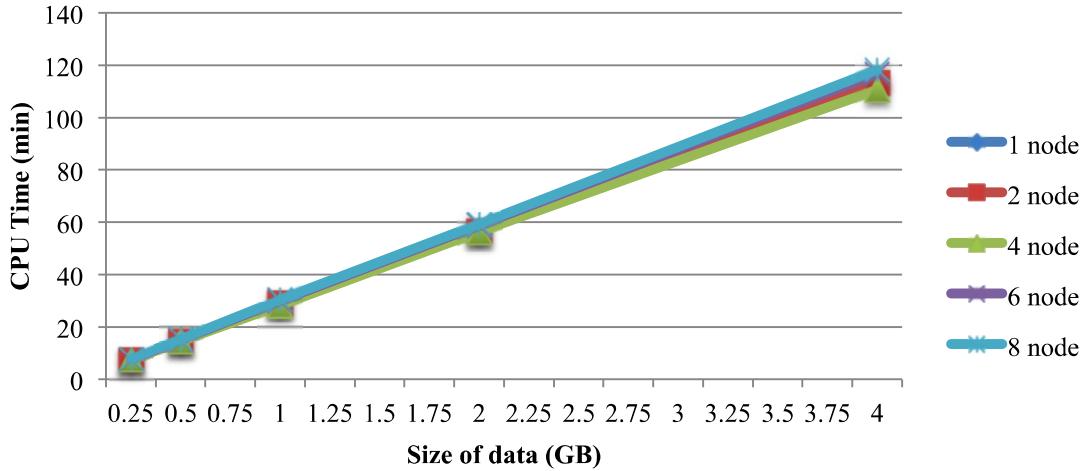


Figure 12: CPU Time and Size of Data

From the above analysis, let us try to find the function to define the relationship between the CPU time and the number of nodes, size of input data. First, let us define some variables. These variables will also be used in the later analysis. Use variable n to define the number of nodes in a cluster. Use variable s to define the size of the input data.

For the CPU time, from the figures and analysis above, we define a function y_1 to represent the relationship. We have known that the CPU time has no relationship with the number of nodes, and is linearly proportional to the size of input data, therefore the function can be written as:

$$y_1 = f(s) = a \cdot s + b$$

In the equation, variables a and b are two coefficients.

5.4 Data Size

In this experiment, we still used M1 large type of node. We have introduced two terminologies before: the total execution time, and the total slot execution time. For this

experiment, we will run the job with different size of input data on multiple clusters with different number of nodes, and try to find the relationship between the time and the data size. In addition, the difference between the total execution time and the total slot execution time could be the time consumed by I/O, and scheduling. We will also have a brief analysis on this time.

5.4.1 Total Execution Time

The total execution time is collected by the program itself. The experiment results are plotted in Figure 13 below.

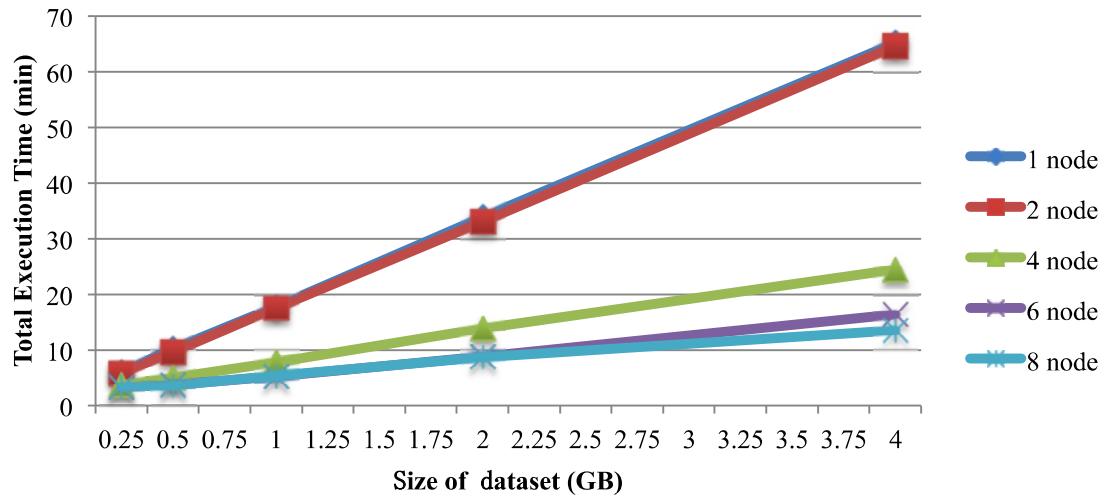


Figure 13: The Total Execution Time and Size of Data

In Figure 13 above, the vertical axis is the total execution time and the unit is minute. The horizontal axis is the size the input data. The legend is the number of nodes. From the above Figure, we can get the following conclusions:

(1) For each node number, the total execution time is linearly related to the size of input data.

This result proves that the MapReduce parallelization does not change the linear relationship between the execution time and the input data size.

(2) The total execution times of 1 node and 2 nodes cluster are the same.

According to our previous analysis, when the number of nodes increases, the total execution time would decrease linearly. However this is not true for the case where the number of nodes increases to 2 from 1. The reason is that when we only had 1 node, this node will serve as both the master node and worker node. Actually it is mainly about worker node because the main computation is consumed by data processing on worker node. However, when we increases the number of nodes to 2, one of them will serve as master node and the other one will serve as the worker node, so actually, there is still only one worker node to process the data. These two cases are almost the same. That is why when the number of nodes increases from 1 to 2 does not have any significant affects on the total execution time.

This is always true for the future situations, that we can see the 1 node cluster and 2 nodes cluster would not have significant differences.

(3) The execution time decreases while the number of nodes increases.

This relationship seems not to be linearly. We will talk more about this in the next section about number of nodes.

5.4.2 Total Slot Execution Time

From the previous section, we have known that the total execution is linearly related to the input data size. However, the total execution time includes many other time consumed in the job, such as the I/O time, and the scheduling time. These times may or may not be proportional to the input data size, or linearly related. Since we have had all the slot times, the differences between them are the time for this part.

First, let us have an overview of the relationship between the total slot execution time and the input data size. It is expected to be linearly related and almost the same as the total execution time relationship. The results are plotted in the following Figure 14:

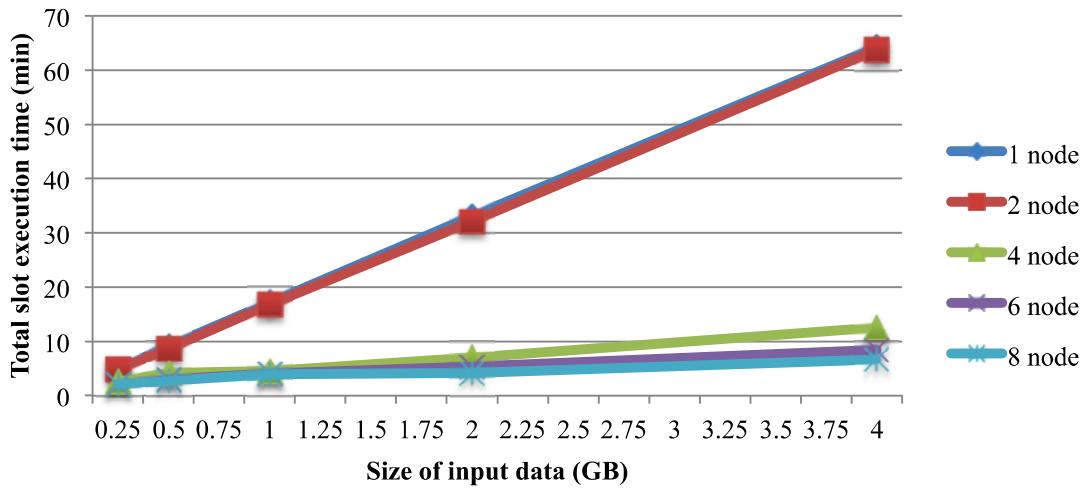


Figure 14: The Total Slot Time and Size of Input Data

In Figure 14 above, the vertical axis is the total slot execution time and the unit is minute. The horizontal axis is the size the input data. The legend is the number of nodes.

From the above Figure, we can get all the same conclusions as the relationship between total execution time and the size of input data in the previous results.

Then let us have a look at the relationship between the differences and the input data size, as Figure 15 shows below.

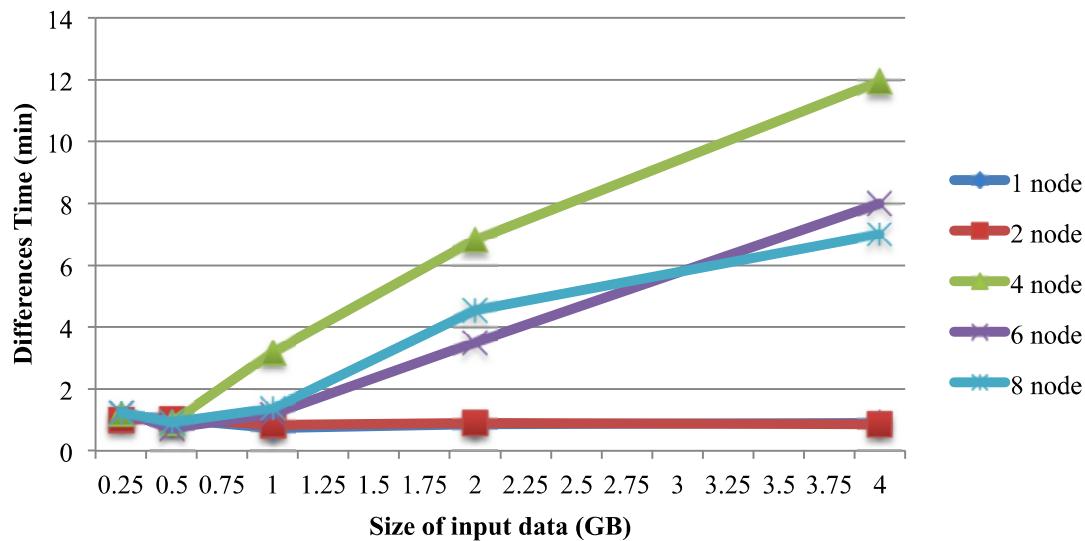


Figure 15: The Difference Time and Size of Input Data

In the above Figure 15, the vertical axis is the difference time between the total execution time and total slot execution time, and the unit is minute. The horizontal axis is the size of input data. The legend is the number of nodes. From the above Figure, we can get the following conclusions:

- (1) For the 1 node and 2 nodes cluster, the difference time stays stable without any changes.**

We have talked about the reason for this phenomenon in the previous section, that the 1 node cluster and 2 nodes cluster do not have significant differences because there is always only 1 worker node in these two clusters. Due to the data locality policy, there would be no data transferred between nodes, which means not too much I/O cost, and the scheduling would be only for a single worker node, which makes it easier and not cost too much time.

(2) The difference time starts to increase after the input data is bigger than 0.5GB, and they are linearly related.

As we mentioned before, this difference time including the I/O time, and the scheduling time, etc. When the data is too small, the I/O cost between nodes might be constant. While when the data size is beyond a value, this I/O cost cannot be ignored any more because it means more data needs to be transferred between nodes and the difference time increases. It is good to see that the relationship is kind of linear, which means the data allocation is well balanced in Hadoop.

From all the above analysis, let us try to find the function to define the relationship between the total execution time and the size of input data. We still use variable s to define the size of the input data. From the figures and analysis above, we define a function $g(s)$ to represent the relationship. We have known that total execution time is linearly proportional to the size of input data; therefore the function can be written as:

$$g(s) = c \cdot s + d$$

In the equation, variables c and d are two coefficients.

5.5 Number of Nodes

In the previous section, we analyzed the relationship between the execution time and the input data size. In this experiment, we will try to analyze the results we get to find the relationship between execution time and the number of nodes in the cluster. We still used M1 large type of node, and run the job with different size of input data on multiple clusters with different number of nodes. We will analyze the relationship between the number of nodes and the total execution time, total slot execution time, and the speedup.

5.5.1 Total Execution Time

In this section we will analyze the relationship between the total execution time and the number of nodes. The results are plotted in the following Figure 16:

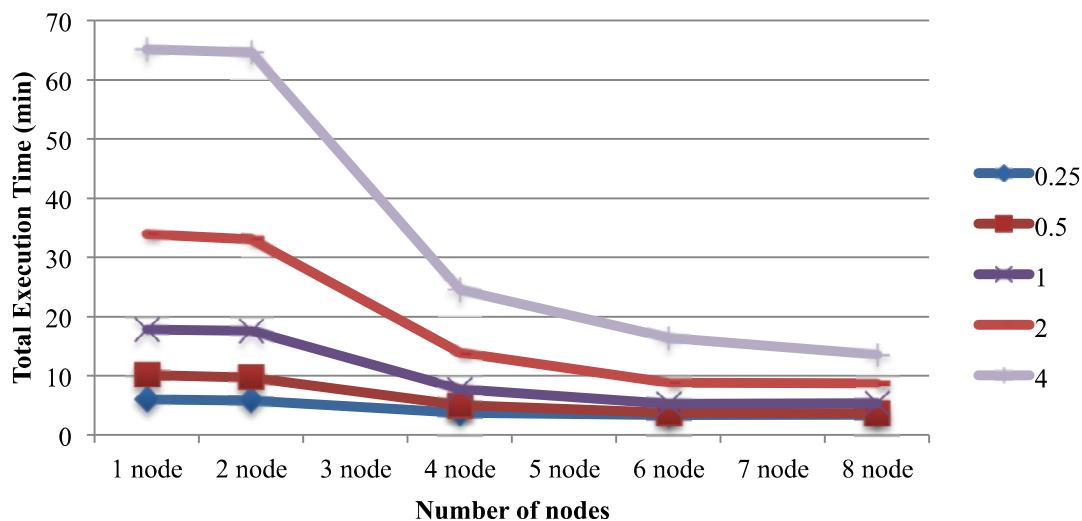


Figure 16: The Total Execution Time and The Number of Nodes

In the above Figure 16, the vertical axis is the total execution time, and the unit is minute. The horizontal axis is the number of nodes. The legend is the input size, and the unit is GB. From the above Figure, we can get the following conclusions:

(1) The total execution time stays the same when the number of node increases from 1 to 2.

We have observed and analyzed this phenomenon in the previous sections.

The reason is that for both 1 node cluster and 2 nodes cluster, there is always only 1 worker node, so the total execution time will not change.

(2) After the number of nodes is beyond 2, the relationship between the total execution time and the number of nodes seems to be inverse proportion.

From the previous chapter of time complexity, we had the conclusion that if we have n nodes in the cluster, the time complexity of the parallel MapReduce Naive Bayes based sentiment analysis algorithm would be:

$$T(n) = \frac{O(c \cdot |V|)}{n}$$

In this equation, c is the number of classes and $|V|$ is the size of feature words. From this equation, we can have a hypothesis that the total execution time should be inverse proportional to the number of nodes. The results in Figure 16 prove that this hypothesis is correct when the number of nodes increases from 2 to 8. And also, the bigger the input data size is, the bigger the size of $|V|$ is, so the graph in Figure 16 is more gently. This experiment results also prove that the previous time complexity analysis is correct.

5.5.2 Total Slot Execution Time

The experiment results for this section are plotted in the following Figure 17.

In Figure 17, the vertical axis is the total slot execution time and the unit is minute.

The horizontal axis is the number of nodes. The legend is the size of input data, and the unit is GB. From the Figure, we can get all the same conclusions as the relationship between total execution time and the number of nodes in the previous discussion.

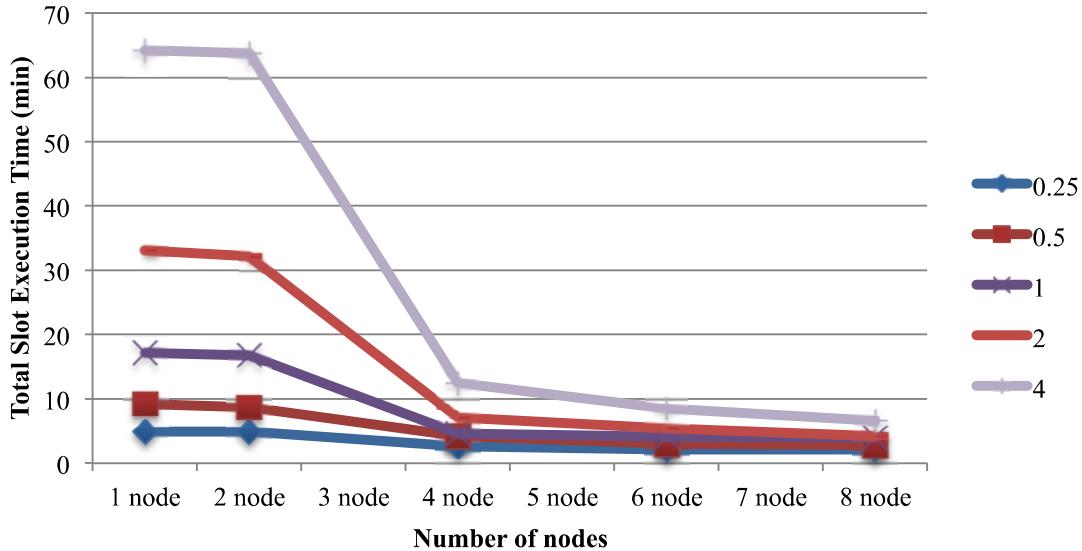


Figure 17: The Total Slot Execution Time and The Number of Nodes

From the above analysis, let us try to find the function to define the relationship between the total execution time and the number of nodes. We still use variable n to define the number of nodes. From the figures and the algorithm analysis above, we define a function $q(n)$ to represent the relationship. We have known that total execution time is

inversely proportional to the size of number of nodes; therefore the function can be written as:

$$q(n) = \frac{e}{n} + f$$

In the equation, variables e and f are two coefficients.

Let us define a new function y_2 to represent the relationship between the total execution time and the number of nodes, size of data. Then we can get:

$$\begin{aligned} y_2 &= g(s) \cdot q(n) = (c \cdot s + d) \cdot \left(\frac{e}{n} + f \right) \\ &= \frac{ce \cdot s}{n} + cf \cdot s + \frac{de}{n} + df \end{aligned}$$

In the equation, variables c , d , e , and f are coefficients. We can use some constant variables to replace them. If the number of nodes is fixed, then the relationship between the total execution time and the size of input data is:

$$y_2 = A \cdot s + B$$

This equation means the total execution time and the size of input data is linearly related. If the size of input data is fixed, then the relationship between the total execution time and the number of nodes is:

$$y_2 = \frac{A'}{n} + B'$$

This equation means the total execution time and the number of nodes is inversely related.

In addition, just as the discussion in the previous section, the relationship between the difference time and the size of input data, here we will have a look at the relationship

between the difference time and the number of node. The result is plotted in the following Figure 18.

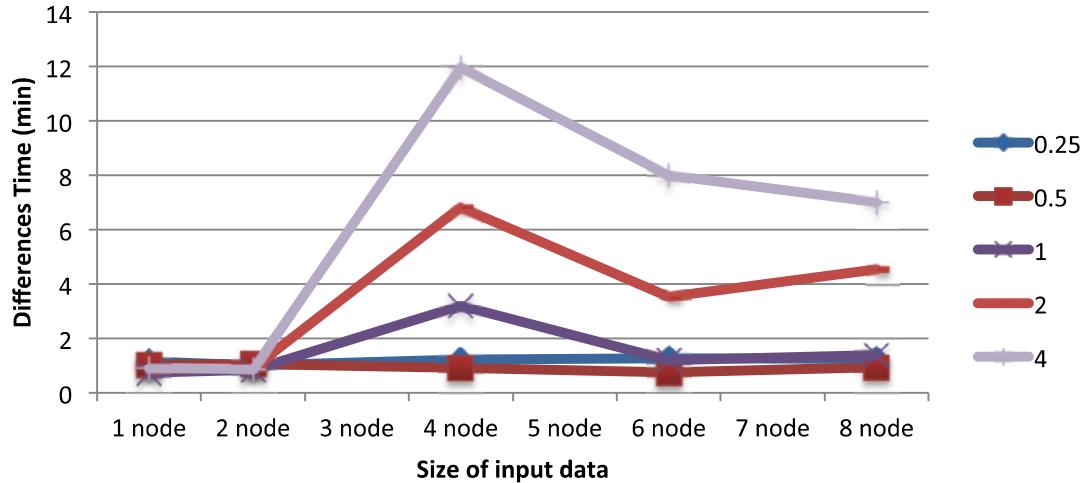


Figure 18: The Difference Time and Size of Input Data

In the above Figure 18, the vertical axis is the difference time between the total execution time and total slot execution time, and the unit is minute. The horizontal axis is the size of input data. The legend is the number of nodes. From the above Figure, we can get the following conclusions:

- (1) For the 1 node and 2 nodes cluster, the difference time stays stable no matter how input size changes.**

We have talked about the reason for this phenomenon in the previous section, that the 1 node cluster and 2 nodes cluster do not have significant differences because there is always only 1 worker node in these two clusters. Due to the data locality policy, there would be no data transferred between nodes, which

means not too much I/O cost, and the scheduling would be only for a single worker node, which makes it easier and not cost too much time.

- (2) When the number of node is 4, it gives the biggest difference time, which means the worst performance. After 4 nodes, the difference time converges to balances status.**

Through the experiments, we figured out that for the input data size is bigger than 1GB, the difference time would have a peak point when the number of node is 4. After that, the difference time will decrease and converges to a constant value. We do not have data for 3 nodes, but the difference time would be high as well when number of node is 3, because when the number of worker nodes increases from 1 (the number of nodes is 2), data stored in HDFS should be duplicated to ensure the stability. The default duplication number of data in Hadoop is 3, so when the number of nodes reaches 4, which means the number of worker node is 3, the data duplication reaches the peak. Data should be transferred in the network 3 times and that will cost a lot of I/O communication. After that, when the number of nodes increases, the data duplication will be distributed over the whole cluster, and finally it will get a balanced status, which is why the difference time would converge to a constant value as the number of nodes increases.

5.5.3 Speedup

The reason why we studied the speedup is to verify the relationships we found before are correct. First, let us define the speedup here. For each job, we have a CPU

time, which can be treated as the ground truth or the sequential time. Also, we have a total execution time, which usually is faster than the CPU time because it is executed in parallel MapReduce model. If we want to know how fast the parallelization can perform, we define a speedup variable like this:

$$Speedup = \frac{CPU\ time}{Total\ execution\ time}$$

Before the analysis, let us have a look at the results we got from our experiments. The results are plotted in the following Figure 19.

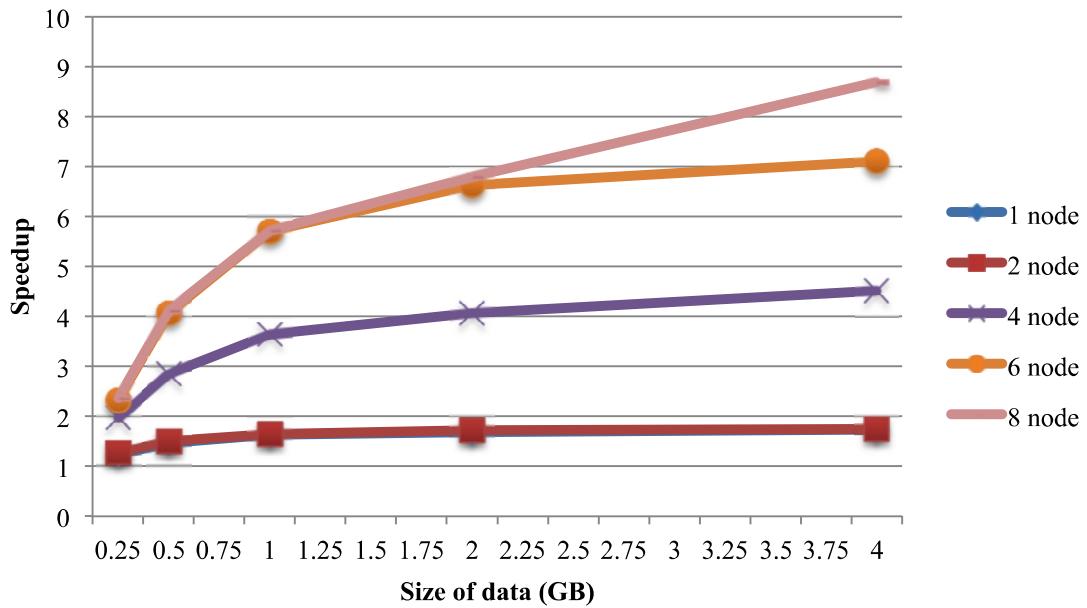


Figure 19: Speedup and Size of Data

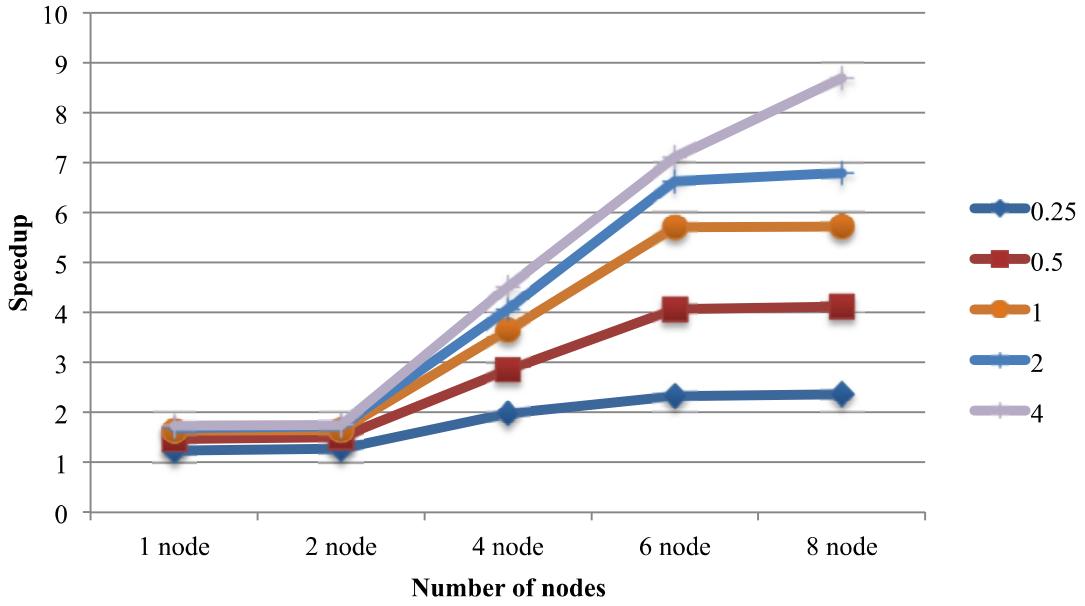


Figure 20: Speedup and The Number of Nodes

We can observe from Figure 20 that the speedup does not change when the number of nodes increases from 1 to 2, which is the same as we analyzed before. While, later as the size of input data and the number of nodes increase, the speedup increases as well. In order to find the relationship, let us use the relationship functions we found before.

We define a function y_3 to represent the relationship between the speedup and the number of nodes, size of data. We have function y_1 to represent the CPU time relationship, and function y_2 to represent the total execution time relationship, then we can get the equation:

$$y_3 = \frac{y_1}{y_2} = \frac{a \cdot s + b}{\frac{ce \cdot s}{n} + cf \cdot s + \frac{de}{n} + df}$$

In the equation, variables a , b , c , d , e , and f are coefficients. We can use some constant variables to replace them. If the number of nodes n is fixed, then the relationship between the speedup and the size of input data s is:

$$y_3 = \frac{A \cdot s + B}{C \cdot s + D} = \frac{A + \frac{B}{s}}{C + \frac{D}{s}}$$

This equation means the speedup is proportional to the size of data s , but not linearly related. And when the size of input s is too large, the speedup would be stable and approach to a constant value. If the size of input data s is fixed, then the relationship between the speedup and the number of nodes n is:

$$y_3 = \frac{A'}{\frac{B'}{n} + C'}$$

This equation means the speedup is also proportional to the number of nodes n , but not linearly related. And when the number of nodes n is too large, the speedup would also be stable and converge to a constant value.

5.6 Hardware Configuration

In addition to the size of input data and the number of nodes in a cluster, the hardware configuration could also affect the performance a MapReduce program. In our experiment, there are two types of nodes, the M1 large, and the M2 4xlarge. Here are the configurations:

Model	vCPU	Memory (GiB)	Storage (GB)
M1 large	2	7.5	2 x 420
M2 4xlarge	8	68.4	4 x 420

We run our program on two 8 nodes cluster with these two kinds of nodes in each cluster. The size of input data is the same as previous experiments.

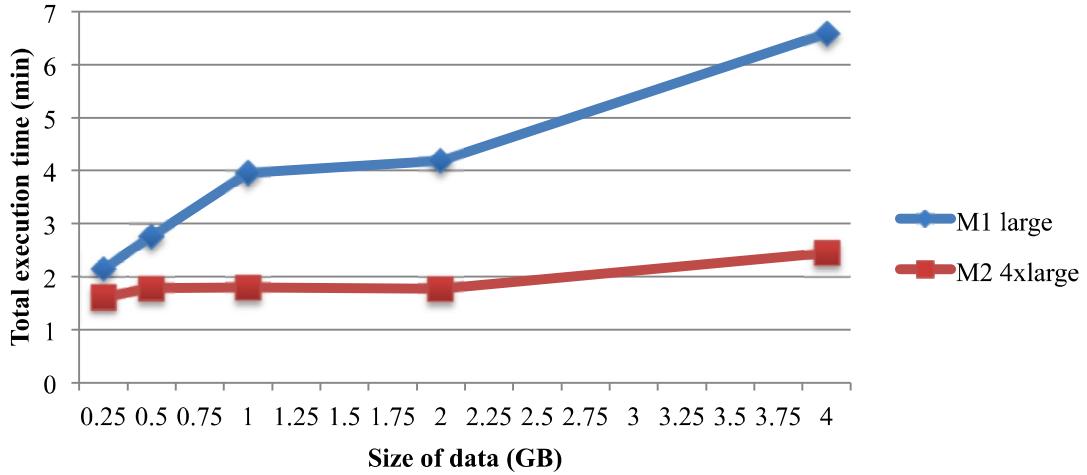


Figure 21: Total Execution Time and Size of Data From Two Clusters

As we can see from the above figure, the larger and more powerful machine can bring some performance improvements. The Amazon Web Service calls the larger node 4xlarge, it means on the whole point the 4xlarge node is 4 times better than the large one. However, from the results we got in Figure 21, the performance improvement is not always exactly 4 times. To some extent, upgrading the hardware can bring performance improvements, but in reality, we still need to make this decision carefully because upgrading hardware of a cluster is hard, and will cost a lot.

5.7 Global Counters

It is hard to share global values in a distributed system. Even if we can, it would give us a lot of problems like the synchronization. However, in MapReduce programming

model, sometimes we need to share some variables or collect some data from each mapper or reducer. Hadoop provides us a mechanism called global counters to collect count data from each node. We have talked about the details about global counters in the previous section. Although global counter is well implemented in Hadoop, it still consumes some resources. We also want to find out how the global variables will affect the performance of our parallel MapReduce Naive Bayes based sentiment analysis algorithm.

As the previous part introduced, there are two user defined global counter in our algorithm, the size of vocabulary, and the unique label number. As comparison, we removed these two global counters and used constant values. The experiment was conducted on 6 nodes cluster with M1 large type of node. The input data is the same as we used before. The total execution time was used to evaluate the performance. The results are plotted in the following Figure 22.

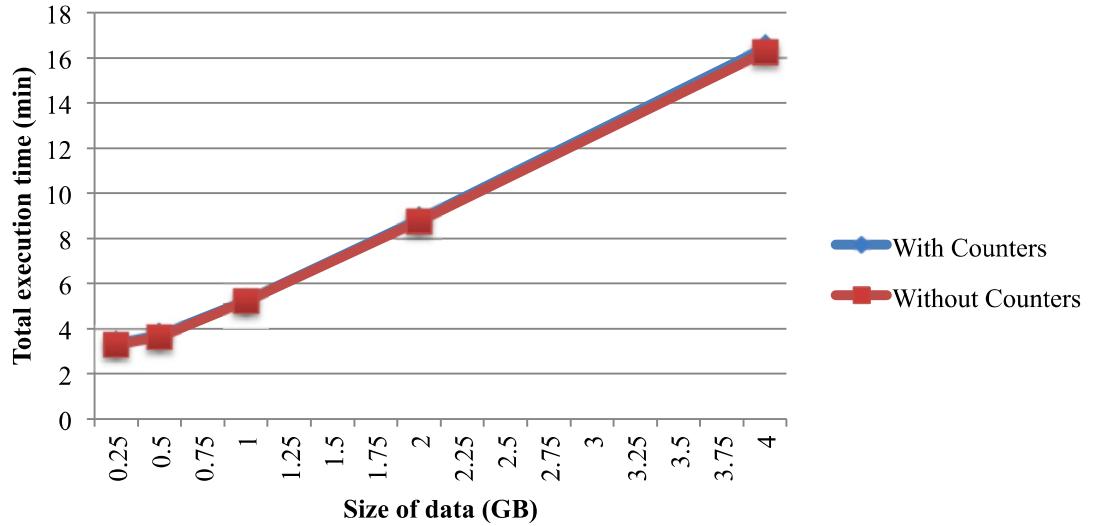


Figure 22: The Total Time Execution and The Size of Data About Counters

As we can see from the Figure 22Figure 22: The Total Time Execution and The Size of Data About Counters, the differences caused by global counters are very small and can be ignored. If we do not use global counters, we need to use other ways to share values, such as writing to a file, which would less efficient than the global counter itself. The conclusion is that use global counters as needed and it would not affect the whole performance of the program very much.

6. CONCLUSION

Text classification is one of the most important areas in text data mining, and sentiment analysis can be performed based on text classification. While the data is exploded today, it becomes a problem to classify large amount of data. The MapReduce programming model solves this problem by providing a mechanism to manage and process data over the cluster.

Naive Bayes is one of the algorithms to perform sentiment analysis. The general idea is to do some word counting and compute some statistics data to build the classifier. A prototype system to find the correlations between geographical Twitter sentiment and stock price using Naive Bayes has been implemented. Through this system, we can collect and analyze the sentiment of tweets based on geographical information, and try to find the correlation between the geographical sentiment and the stock prices of companies. The preliminary results show that for the Starbucks Corporation, there is a weak correlation between them, and among the five the states the Illinois State has the strongest correlation, while the Texas State has the weakest correlation.

The Naive Bayes sentiment analysis algorithm has also been proved to be suitable for MapReduce programming model. Based on Hadoop framework, we implemented the parallel Naive Bayes based sentiment analysis algorithm in MapReduce model, and through a series of experiments and analysis, the MapReduce implementation proves to have scalability, and we also provide a prediction model to predict and analyze the performance of parallel Naive Bayes MapReduce program. This prediction model can also be used for other MapReduce programs.

7. FUTURE WORK

Although we have made a step towards the analysis and prediction of Naive Bayes MapReduce implementation, there is still much more work that needs more and deeper research.

- (1) The collected sample data is limited. The streaming data may be big enough, however, less than 5% of them have geographical information available. The amount of tweets finally flowed into our system are very small. A method to predict or deduct the location of a tweet based on the tweet's information and the user's information should be found in the future.
- (2) For the parallel Naive Bayes MapReduce implementation in this thesis, it is not only can be applied to text classification or sentiment analysis. It can also be extended to other areas like image processing, or gene analysis etc.
- (3) We have concluded the performance prediction model related to the size of input data and the number of nodes in a cluster, we also have a certain understanding of the relationship between the performance and the global counters, and the hardware configuration. For the latter, we still need more data and experiments to have more details, such as how many global counters we should use for a certain use case, or under what situation that upgrading the hardware is worth the improvements it brings.
- (4) When the MapReduce framework is dealing with some algorithms that have a high time complexity in the data mining area, the consumption of resource is increased since the exchange of data will bring a lot of I/O transferring and

network communication during the iteration. Therefore, how to reduce and optimize the data exchange in the MapReduce framework is also an issue that needs to be solved.

8. BIBLIOGRAPHY

- [1] Manyika, James, et al. "Big data: The next frontier for innovation, competition, and productivity." (2011). A
- [2] Dean, Jeffrey, and Sanjay Ghemawat. "MapReduce: simplified data processing on large clusters." Communications of the ACM 51.1 (2008): 107-113.
- [3] Gropp, William, et al. "A high-performance, portable implementation of the MPI message passing interface standard." Parallel computing 22.6 (1996): 789-828.
- [4] Hadoop, Apache. "Hadoop." [2009-03-06]. <http://hadoop.apache.org> (2009).
- [5] Temple, Krystal. "What Happens in an Internet Minute?." Inside Scoop (2012). A
- [6] Zaharia, Matei, et al. "Spark: cluster computing with working sets." Proceedings of the 2nd USENIX conference on Hot topics in cloud computing. 2010.
- [7] Zaharia, Matei, et al. "Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing." Proceedings of the 9th USENIX conference on Networked Systems Design and Implementation. USENIX Association, 2012.
- [8] Chaiken, Ronnie, et al. "SCOPE: easy and efficient parallel processing of massive data sets." Proceedings of the VLDB Endowment 1.2 (2008): 1265-1276.
- [9] Isard, Michael, et al. "Dryad: distributed data-parallel programs from sequential building blocks." ACM SIGOPS Operating Systems Review. Vol. 41. No. 3. ACM, 2007.
- [10] Hindman, Benjamin, et al. "Mesos: A Platform for Fine-Grained Resource Sharing in the Data Center." NSDI. Vol. 11. 2011.
- [11] Bu, Yingyi, et al. "HaLoop: Efficient iterative data processing on large clusters." Proceedings of the VLDB Endowment 3.1-2 (2010): 285-296.
- [12] Malewicz, Grzegorz, et al. "Pregel: a system for large-scale graph processing." Proceedings of the 2010 ACM SIGMOD International Conference on Management of data. ACM, 2010.
- [13] Marz, Nathan. "Storm-distributed and fault-tolerant realtime computation." (2013).
- [14] [1] MNIST handwritten digits image recognition dataset. Available via [www.
http://yann.lecun.com/exdb/mnist/](http://yann.lecun.com/exdb/mnist/)

- [15] Shvachko, Konstantin, et al. "The hadoop distributed file system." Mass Storage Systems and Technologies (MSST), 2010 IEEE 26th Symposium on. IEEE, 2010.
- [16] Pang, Bo, Lillian Lee, and Shivakumar Vaithyanathan. "Thumbs up?: sentiment classification using machine learning techniques." Proceedings of the ACL-02 conference on Empirical methods in natural language processing- Volume 10. Association for Computational Linguistics, 2002.
- [17] Taboada, Maite, et al. "Lexicon-based methods for sentiment analysis." Computational linguistics 37.2 (2011): 267-307.
- [18] McCallum, Andrew, and Kamal Nigam. "A comparison of event models for naive bayes text classification." AAAI-98 workshop on learning for text categorization. Vol. 752. 1998.
- [19] Yousef, Ahmed Hassan, Walaa Medhat, and Hoda Korashy Mohamed. "Sentiment Analysis Algorithms and Applications: A Survey." (2014).
- [20] Mullen, Tony, and Nigel Collier. "Sentiment Analysis using Support Vector Machines with Diverse Information Sources." EMNLP. Vol. 4. 2004.
- [21] Dobro, Alin. "Decision Tree Classification." Encyclopedia of Database Systems. Springer US, 2009. 765-769.
- [22] Kim, Susanna. "Twitter's IPO Filing Shows 215 Million Monthly Active Users." ABC News 3 (2013).
- [23] Guide, Developer. "Amazon Elastic MapReduce." (2010).
- [24] Bollen, Johan, Huina Mao, and Xiaojun Zeng. "Twitter mood predicts the stock market." Journal of Computational Science 2.1 (2011): 1-8.
- [25] E. Gilbert, K. Karahalios, Widespread worry and the stock market, in: Fourth International AAAI Conference on Weblogs and Social Media, Washington, DC, 2010, pp. 58–65,
<http://www.aaai.org/ocs/index.php/ICWSM/ICWSM10/paper/download/1513/1833>.
- [26] Bollen, Johan, Huina Mao, and Alberto Pepe. "Modeling public mood and emotion: Twitter sentiment and socio-economic phenomena." ICWSM. 2011.
- [27] Kulshrestha, Juhi, et al. "Geographic Dissection of the Twitter Network." ICWSM. 2012.
- [28] Mislove, Alan, et al. "Understanding the Demographics of Twitter Users." ICWSM 11 (2011): 5th.
- [29] Google Finance: <https://www.google.com/finance>
- [30] Rao, Tushar, and Saket Srivastava. "Analyzing stock market movements using twitter sentiment analysis." Proceedings of the 2012 International Conference on Advances in Social Networks Analysis and Mining (ASONAM 2012). IEEE Computer Society, 2012.